

SSH Shell Attacks

ANDREA BOTTICELLA*, Politecnico di Torino, Italy

ELIA INNOCENTI*, Politecnico di Torino, Italy

RENATO MIGNONE*, Politecnico di Torino, Italy

SIMONE ROMANO*, Politecnico di Torino, Italy

This paper introduces a comprehensive machine learning approach to analyze SSH shell attack sessions, leveraging both supervised and unsupervised learning techniques. Using a dataset of 230,000 unique Unix shell attack sessions, the framework aims to classify attacker tactics based on the MITRE ATT&CK framework and uncover latent patterns through clustering.

The key contributions of this work are:

- Development of a robust pre-processing pipeline to analyze temporal trends, extract numerical features, and evaluate intent distributions from large-scale SSH attack session data.
- Implementation of supervised classification models to accurately predict multiple attacker tactics, supported by hyperparameter tuning and feature engineering for enhanced performance.
- Application of unsupervised clustering techniques to uncover hidden patterns in attack behaviors, leveraging visualization tools and cluster analysis for fine-grained categorization.
- Exploration of advanced language models, such as BERT and Doc2Vec, for representation learning and fine-tuning to improve intent classification and session interpretation.

CCS Concepts: • **Computing methodologies** → **Supervised learning by classification; Unsupervised learning; Natural language processing; Machine learning; Machine learning approaches**; • **Security and privacy** → Intrusion detection systems.

Additional Key Words and Phrases: Machine learning, supervised learning, unsupervised learning, language models, text classification, clustering, intent classification, SSH shell attacks, security log analysis

CONTENTS

Abstract	1
Contents	1
1 INTRODUCTION	1
2 BACKGROUND	2
3 DATA EXPLORATION AND PRE-PROCESSING	3
4 SUPERVISED LEARNING - CLASSIFICATION	8
5 UNSUPERVISED LEARNING - CLUSTERING	15
6 LANGUAGE MODEL EXPLORATION	17
7 CONCLUSION	20

1 INTRODUCTION

1.1 Motivation

Security logs play a crucial role in understanding and mitigating cyber attacks, particularly in the domain of network and system security. With the increasing sophistication of cyber threats, analyzing and interpreting security logs has become paramount for detecting, preventing, and responding to potential security breaches.

*The authors collaborated closely in developing this project.

Authors' Contact Information: Andrea Botticella, andrea.botticella@studenti.polito.it, Politecnico di Torino, Turin, Italy; Elia Innocenti, elia.innocenti@studenti.polito.it, Politecnico di Torino, Turin, Italy; Renato Mignone, renato.mignone@studenti.polito.it, Politecnico di Torino, Turin, Italy; Simone Romano, simone.romano2@studenti.polito.it, Politecnico di Torino, Turin, Italy.

Unix shell attacks, especially those executed through SSH, represent a significant vector for potential system compromises.

The complexity of security log analysis stems from several key challenges:

- Logs are often unstructured and contain ambiguous or malformed text.
- Manual parsing and interpretation of logs is time-consuming and error-prone.
- The sheer volume of log data makes comprehensive manual review impractical.

These challenges underscore the need for automated, intelligent approaches to log analysis that can efficiently extract meaningful insights and identify potential security threats.

1.2 Objective

The primary objective of this research is to develop and evaluate machine learning techniques for automatic analysis and classification of SSH shell attack logs. Specifically, we aim to:

- Automate the process of log analysis and intent classification.
- Provide security professionals with insights into attack strategies.

The significance of this research lies in its potential to enhance cybersecurity threat detection and response capabilities by transforming complex, unstructured log data into actionable intelligence.

2 BACKGROUND

2.1 Security Logs and Attack Analysis

In the context of SSH shell attacks, logs document the sequence of commands executed during a malicious session, enabling security researchers to analyze attacker behaviors, techniques, and potential system impacts. However, the manual analysis of these logs is challenging due to their volume, complexity, and often non-standard formatting.

2.2 MITRE ATT&CK Framework

The MITRE ATT&CK (Adversarial Tactics, Techniques, and Common Knowledge) framework provides a comprehensive knowledge base of adversary tactics and techniques observed in real-world cyber attacks. This framework serves as a standardized methodology for understanding and categorizing attack strategies.

For our research, we focus on seven key intents derived from the MITRE ATT&CK framework:

- **Persistence:** Techniques used by adversaries to maintain system access across restarts or credential changes.
- **Discovery:** Methods for gathering information about the target system and network environment.
- **Defense Evasion:** Strategies to avoid detection by security mechanisms.
- **Execution:** Techniques for running malicious code on target systems.
- **Impact:** Actions aimed at manipulating, interrupting, or destroying systems and data.
- **Other:** Less common tactics including Reconnaissance, Resource Development, Initial Access, etc.
- **Harmless:** Non-malicious code or actions.

2.3 Research Approach

Our research employs a multi-faceted approach to SSH shell attack log analysis:

- Explore and preprocess a large dataset of Unix shell attack sessions.
- Apply supervised learning techniques to classify attack tactics based on session characteristics.
- Utilize unsupervised learning methods to discover patterns and clusters in attack sessions.
- Investigate the potential of advanced language models in understanding and categorizing attack intents.

3 DATA EXPLORATION AND PRE-PROCESSING

3.1 Introduction

This section outlines the steps taken to explore and preprocess the dataset used in this study. The primary objective is to understand the data characteristics, identify patterns, and prepare the data for further analysis. We focus on temporal trends, intent distributions, and textual features.

3.2 Dataset Overview

The dataset consists of logs from SSH attacks. Each entry includes the following features:

- **session_id**: An integer representing the unique identifier for each session.
- **full_session**: A string containing the full sequence of commands executed during the attack session.
- **first_timestamp**: The timestamp indicating the start of the session.
- **Set_Fingerprint**: A set of strings (labels) describing the nature of the attack.

Below is an example of the dataset structure:

	session_id	full_session	first_timestamp	Set_Fingerprint
0	0	enable; system; shell; sh; cat /proc/mounts; /...	2019-06-04 09:45:11.151186+00:00	[Defense Evasion, Discovery]
1	1	enable; system; shell; sh; cat /proc/mounts; /...	2019-06-04 09:45:50.396610+00:00	[Defense Evasion, Discovery]
2	2	enable; system; shell; sh; cat /proc/mounts; /...	2019-06-04 09:54:41.863315+00:00	[Defense Evasion, Discovery]
...
233034	233046	cat /proc/cpuinfo grep name wc -l; echo -e...	2020-02-29 23:59:22.199490+00:00	[Discovery, Persistence]

233035 rows x 4 columns

Fig. 1. Dataset

3.3 Dataset Preparation and Feature Extraction

The dataset used in this study was provided as a Parquet file.

To prepare the dataset for analysis and ensure its quality, several preprocessing steps were undertaken, described in detail below:

- (1) **Decoding the full_session column**: The full_session column contained 90026 encoded shell scripts, making the raw data difficult to interpret. A decoding process was applied to these entries, converting them into a human-readable format.
- (2) **Formatting the first_timestamp column**: The first_timestamp column was checked for consistency and converted into a standard datetime format.
- (3) **Splitting the full_session column into lists of commands**: The full_session column contained entire attack sessions represented as single strings. Since we wanted each word / command to be a feature of our models, it was necessary to break it down into individual commands or keywords for more granular analysis. This was achieved through a multi-step splitting process using specific delimiters:
 - The semicolon (;) was used to separate distinct commands within a session.
 - The pipe (|) was used to divide concatenated commands or pipelines.
 - Whitespace () was used to further split commands into individual words or arguments.

This process allowed the identification of specific actions and parameters, facilitating a detailed analysis of the attack strategies.

- (4) **Cleaning individual commands:** Each command or keyword extracted from the `full_session` column was processed to remove unnecessary or undesired elements. Regular expressions were used to strip unwanted symbols, qualifiers, and variable assignments.
- (5) **Handling missing values:** To ensure the integrity and completeness of the dataset, a thorough check for missing values was performed across all columns.

```
enable ; system ; shell ; sh ; cat /proc/mounts ; /bin/busybox SAEMW ; cd /dev/
shm ; cat .s || cp /bin/echo .s ; /bin/busybox SAEMW ; tftp ; wget ; /bin/
busybox SAEMW ; dd bs=52 count=1 if=.s || cat .s || while read i ; do echo
$i ; done < .s ; /bin/busybox SAEMW ; rm .s ; exit ;

[enable, system, shell, sh, cat, /proc/mounts, /bin/busybox, SAEMW, cd, /dev/shm
, cat, .s, cp, /bin/echo, .s, /bin/busybox, SAEMW, tftp, wget, /bin/busybox
, SAEMW, dd, bs, count, if, cat, .s, while, read, i, do, echo, i, done, .s,
/bin/busybox, SAEMW, rm, .s, exit]
```

Listing 1. Dataset Processing

These preprocessing steps ensured that the dataset was well-structured, clean, standardized and ready for advanced analysis and machine learning tasks. By addressing issues such as encoding, formatting, and missing data, the preprocessing phase established a robust foundation for the study.

3.4 Temporal Analysis

The temporal analysis examines when the attacks were performed and the intent distribution over time.




Fig. 2. Temporal Series of SSH Attacks

Our dataset spans from June 2019 to March 2020, revealing a notable concentration of attacks between mid-October 2019 and January 2020.

Further analysis was conducted on the distribution of attacks/sessions across various temporal dimensions (months, days of the week, etc.), but no significant patterns were identified.

3.5 Common Words Analysis

The analysis of common words within the attack sessions provides insights into frequently utilized commands and keywords. These elements are pivotal for understanding attacker behavior and informing model training for detecting attack intents.

To visualize and analyze the most frequent words, two approaches were used:

- **Word Cloud Visualization:** A word cloud (Figure 3a) was generated to represent the most common words in the dataset. Commands like `tmp`, `grep`, `var`, and `cat` dominate the word cloud, highlighting their frequent usage in attack sessions.
- **Top 20 Most Common Words Bar Plot:** A bar plot (Figure 3b) showcases the top 20 most frequent words, ordered by their frequency. This plot provides more quantitative detail, revealing the high occurrence of commands such as `tmp` (over 2 million occurrences) and `var` (nearly 2 million occurrences). Other significant commands include `grep`, `cat`, `echo`, and `cp`, each surpassing 1 million occurrences. These commands are commonly associated with file manipulation, string searching, and process information, which are typical in malicious activities.

The combination of these visualizations supports the identification of commonly used commands and their potential roles in attack sessions.

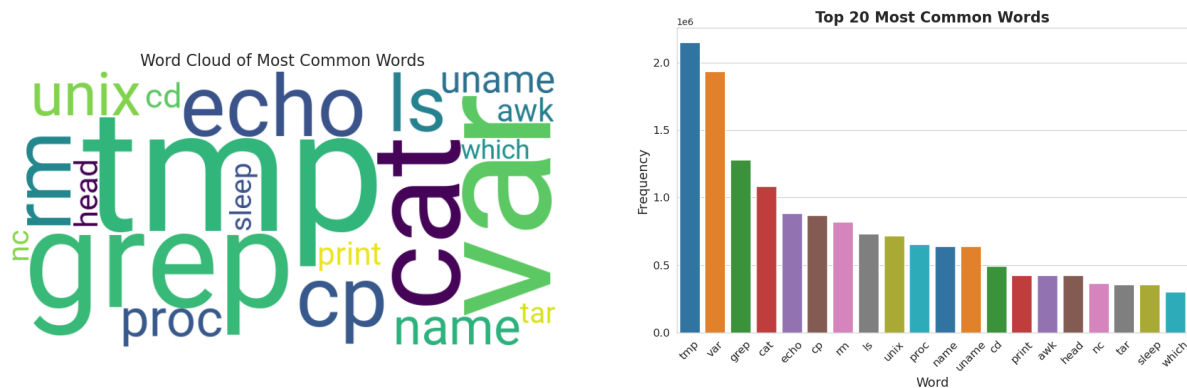


Fig. 3. (a) Word Cloud of Most Common Words and (b) Top 20 Most Common Words

3.6 Intent Analysis

This section presents an analysis of different intents identified in attack sessions and their relationships. Our analysis reveals distinct patterns in both the frequency of individual intents and their co-occurrence relationships, providing valuable insights into attacker behavior patterns.

Figure 4a illustrates the frequency distribution of different intents in our dataset. The analysis reveals that Discovery and Persistence are the most prevalent intents, with approximately 200,000 occurrences each. This suggests that attackers primarily focus on reconnaissance activities and establishing long-term presence in the

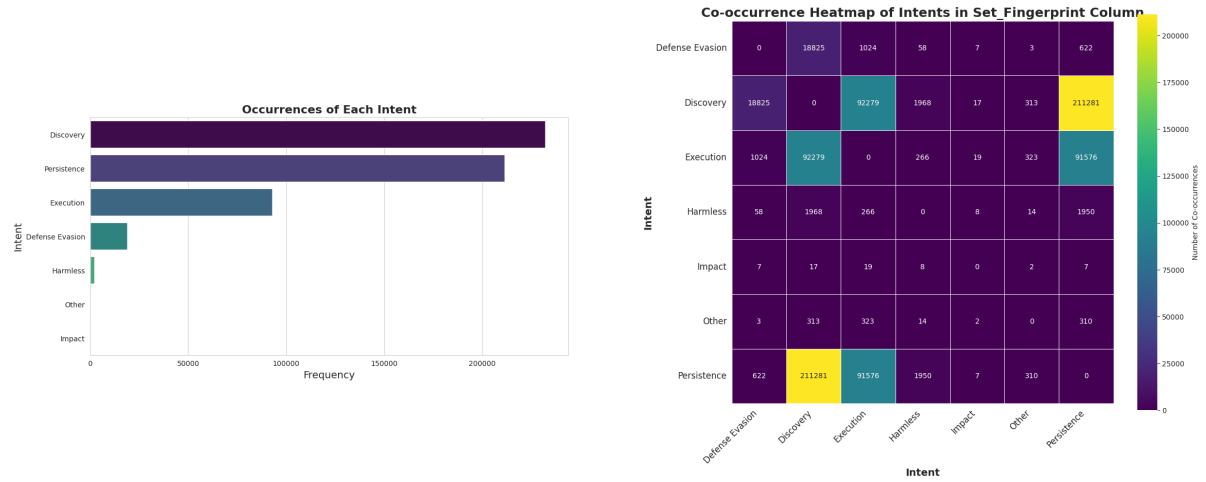


Fig. 4. (a) Occurrences of each Intent and (b) Co-Occurrence Heatmap of Intents

targeted systems. Execution intent appears as the third most common, with roughly 100,000 occurrences, indicating a significant number of attempts to run malicious code. Defense Evasion shows notably fewer occurrences (around 20,000), while Harmless, Other, and Impact intents are relatively rare in the dataset.

The co-occurrence heatmap (Figure 4b) reveals significant patterns in how different intents interact within attack sessions. Several key observations emerge:

- The strongest co-occurrence relationship exists between Discovery and Persistence (211,281 co-occurrences), suggesting that attackers often combine reconnaissance activities with attempts to maintain system access.
- Execution shows strong correlations with both Discovery (92,279) and Persistence (91,576), indicating that malicious code execution frequently accompanies both system discovery and persistence establishment attempts.
- Defense Evasion exhibits moderate co-occurrence with Discovery (18,825) and minimal correlation with other intents, suggesting that evasion techniques are primarily employed during reconnaissance phases.
- Harmless, Impact, and Other intents show minimal co-occurrence with other categories (with values mostly under 20 co-occurrences). This isolation pattern is partially explained by their low frequency in the dataset: Harmless with 2,206 occurrences, Impact with only 27 occurrences, and Other with 327 occurrences. These numbers represent a small fraction of the total recorded intents, naturally limiting their potential for co-occurrence with other intent types.

These findings highlight common attack patterns where adversaries typically begin with discovery operations, followed by persistence establishment and execution of malicious code. This understanding can inform the development of detection strategies and defensive measures, particularly focusing on the most common intent combinations.

3.7 Session Analysis

The session analysis aims to understand the structural characteristics of the attack sessions in the dataset. This is accomplished through two visualizations: the Empirical Cumulative Distribution Function (ECDF) plots and the distribution of the number of words per session.

The ECDF plots in Figure 5a provide insights into the distribution of the number of characters and words across all sessions. The Empirical Cumulative Distribution Function (ECDF) is a statistical tool that represents the proportion of observations below a certain value, offering an intuitive way to understand data distributions.

- The ECDF for characters indicates that the majority of sessions have fewer than 10^3 characters, with a sharp increase between 10^2 and 10^3 characters. This suggests a high concentration of relatively short sessions.
- The ECDF for words shows a similar trend, with most sessions containing fewer than 10^2 words. The distribution reflects the concise nature of many attack sessions, potentially focusing on executing a limited number of commands.




Fig. 5. (a) ECDF for Characters and Words in Sessions and (b) Distribution of the Number of Words per Session

Figure 5b illustrates the density distribution of the number of words per session. Key observations include:

- The distribution is bimodal, with two distinct peaks. This likely reflects different categories of sessions, such as exploratory attacks with a larger number of commands versus simpler, targeted attacks with fewer commands.
- The majority of sessions have fewer than 100 words, reinforcing the compactness observed in the ECDF plots.
- Outliers with higher word counts likely represent complex sessions, potentially involving multiple stages or more sophisticated attack strategies.

The relevance of these plots lies in their ability to provide a foundational understanding of session structure, which is critical for feature engineering and model development. Shorter sessions may correspond to quick, automated attacks, while longer sessions might represent manual or multi-step intrusions. These insights can inform the design of models by emphasizing features tailored to the varying lengths and complexities of sessions, ultimately improving detection accuracy.

3.8 Text Representation

To enable further analysis of the dataset, the initial textual data was transformed into numerical representations. This step is essential for applying machine learning techniques, as numerical formats are required for model training and evaluation. Two widely used text representation techniques were employed:

- **Bag of Words (BoW):** This technique represents each session as a vector where each dimension corresponds to a specific word, and the value represents the frequency of that word in the session. While simple and interpretable, BoW does not capture the importance or uniqueness of words across the dataset.
- **Term Frequency-Inverse Document Frequency (TF-IDF):** This method extends BoW by weighting the word frequencies based on their inverse document frequency, thereby emphasizing words that are

important within a session but less frequent across all sessions. This approach helps reduce the impact of common but non-informative words.

Figure 6 and Figure 7 illustrate the transformed datasets using the BoW and TF-IDF techniques, respectively. The TF-IDF representation was selected for further analysis due to its ability to capture more meaningful patterns in the data by accounting for word relevance.

The use of the TF-IDF representation ensures that subsequent analyses focus on the most distinctive features of each session, enhancing the effectiveness of machine learning models in identifying patterns and predicting intents.

	Set_Fingerprint	enable	system	sh	cat	proc	mounts	bin	busybox	cd	_	top	sleep	tar	passwd	nc	mkdir	unix	mv	nohup	tss
0	[Defense Evasion, Discovery]	0.000000	0.000000	0.001378	0.005652	0.001109	0.001295	0.00042	0.0	0.007463	...	0.0	0.002448	0.000000	0.000000	0.0	0.000000	0.0	0.000000	0.000559	0.000000
1	[Defense Evasion, Discovery]	0.000000	0.000000	0.001378	0.005652	0.001109	0.001295	0.00042	0.0	0.007463	...	0.0	0.002448	0.000000	0.000000	0.0	0.000000	0.0	0.000000	0.000559	0.000000
2	[Defense Evasion, Discovery]	0.000000	0.000000	0.001378	0.005652	0.001109	0.001295	0.00042	0.0	0.007463	...	0.0	0.002448	0.000000	0.000000	0.0	0.000000	0.0	0.000000	0.000559	0.000000
...
233034	[Discovery, Persistence]	0.002165	0.002154	0.000000	0.000000	0.000000	0.001727	0.00000	0.0	0.000000	...	0.0	0.000000	0.000776	0.002152	0.0	0.002166	0.0	0.001611	0.000000	0.001432

233035 rows × 46 columns

Fig. 6. Bag-of-Words

	Set_Fingerprint	enable	system	sh	cat	proc	mounts	bin	busybox	cd	_	top	sleep	tar	passwd	nc	mkdir	unix	mv	nohup	tss
0	[Defense Evasion, Discovery]	0.000000	0.000000	0.656004	0.121018	0.52813	0.111565	0.064612	0.0	0.132227	...	0.0	0.131995	0.000000	0.000000	0.0	0.000000	0.0	0.000000	0.130346	0.000000
1	[Defense Evasion, Discovery]	0.000000	0.000000	0.656004	0.121018	0.52813	0.111565	0.064612	0.0	0.132227	...	0.0	0.131995	0.000000	0.000000	0.0	0.000000	0.0	0.000000	0.130346	0.000000
2	[Defense Evasion, Discovery]	0.000000	0.000000	0.656004	0.121018	0.52813	0.111565	0.064612	0.0	0.132227	...	0.0	0.131995	0.000000	0.000000	0.0	0.000000	0.0	0.000000	0.130346	0.000000
...
233034	[Discovery, Persistence]	0.136394	0.068782	0.000000	0.000000	0.00000	0.252226	0.000000	0.0	0.000000	...	0.0	0.000000	0.340026	0.068061	0.0	0.204608	0.0	0.477045	0.000000	0.06821

233035 rows × 46 columns

Fig. 7. TF-IDF

4 SUPERVISED LEARNING - CLASSIFICATION

4.1 Introduction

The supervised learning experiment in this project aimed to classify attack sessions into various intent categories derived from the ‘Set_Fingerprint’ column of the dataset. This section explores the use of Random Forest and Support Vector Machines (SVM) as the primary models. The analysis focuses on how these models handle multi-label classification and evaluates their performance using metrics such as weighted F1-scores and confusion matrices.

This section outlines the model training and evaluation processes, and a detailed discussion of the results obtained. Each model’s strengths and weaknesses are analyzed, providing insights into their application to multi-label classification problems.

4.2 Data Preprocessing

To prepare the data for supervised learning:

- (1) **Dataset loading:** To perform the supervised learning process, we decided to start from a TF-IDF representation of the dataset, thanks to which it was possible to assign a weight to each word of the SSH session attack.
- (2) **Encoding Intents:** The ‘Set_Fingerprint’ column was encoded into multi-label binary format using the MultiLabelBinarizer. Each intent was represented as a binary vector, allowing for simultaneous prediction of multiple labels.
- (3) **Splitting the Data:** The dataset was divided into training (70%), validation (20%), and testing (10%) subsets. Stratified splitting was employed to ensure that all intents were well-represented in each subset.

These steps ensured the dataset was clean, balanced, and ready for supervised learning.

4.3 Model Training

Three models were trained and evaluated using their default configurations to establish baseline performance:

1. *Random Forest*

The Random Forest model was trained with default parameters, including 100 estimators and unlimited maximum depth. This initial training provided insights into potential overfitting or underfitting issues and served as a benchmark for subsequent tuning.

2. *Support Vector Machines (SVM)*

SVM was initially trained with default settings using a linear kernel and a regularization parameter $C = 1$. The performance was evaluated to assess the model's ability to handle multi-label classification tasks with linearly separable data.

3. *Logistic Regression*

In order to have another view of the analysis, we performed the training with the logistic regression model. While the model performed well, the results were not as significant as those of RF and SVM, so they will not be discussed in detail in this section (see Appendix).

4.4 Evaluation Metrics

The models were evaluated using the following metrics:

- **Accuracy, Precision, Recall**
- **Confusion Matrices:** Provided insight into true positives, false positives, false negatives, and true negatives for each intent.
- **Weighted F1-Scores:** Measured the harmonic mean of precision and recall, with weights proportional to class support.

This evaluation allowed for the identification of baseline performance, highlighting potential areas for improvement through hyperparameter tuning.

4.5 Hyperparameter Tuning

To optimize each model, hyperparameter tuning was performed using a grid search approach. This process aimed to improve performance and address issues of overfitting or underfitting observed in the baseline models:

1. *Random Forest*

The grid search explored combinations of the number of estimators (50, 100, and 150) and maximum depth (unlimited, 50, and 100). The best-performing configuration was selected based on weighted F1-scores.

2. *Support Vector Machines (SVM)*

For SVM, the grid search varied the regularization parameter C (0.01, 0.1, 1, and 10) and the kernel type (linear and RBF). Additional tuning for the RBF kernel included the gamma parameter (scale and auto).

4.6 Results and Observations

4.6.1 *Random Forest*

- *Performance Overview with base model*

The Random Forest base model showed robust performance, achieving high weighted F1-scores across all intents. However, minor overfitting was observed in intents with smaller sample sizes. Figure ?? illustrates the confusion matrix for the base model, showing excellent precision and recall for major intents like *Persistence* and *Defense Evasion*.

- *Hyperparameter Analysis*

Hyperparameter tuning improved the Random Forest model's performance, particularly for minority intents. The optimized model reduced overfitting by limiting tree depth and increasing the number of estimators. Figure ?? highlights the improvement in weighted F1-scores across different configurations.

- *Comparative Analysis of Baseline and Optimized Models*

Comparing the base and optimized models revealed significant improvements in precision and recall for minority intents. The optimized model maintained high performance for major intents while addressing issues of overfitting. Figure ?? presents the confusion matrix for the optimized Random Forest model.

4.6.2 Support Vector Machines (SVM)

- *Performance Overview with base model*

The SVM base model performed well for linearly separable data, achieving high precision and recall for intents with large sample sizes. However, it struggled with minority intents due to its sensitivity to class imbalances. Figure ?? illustrates the confusion matrix for the base model.

- *Hyperparameter Analysis*

Grid search tuning improved SVM's performance, especially for the RBF kernel. Increasing the regularization parameter C and fine-tuning the gamma parameter enhanced the model's ability to classify minority intents. Figure ?? shows the weighted F1-scores for different hyperparameter combinations.

- *Comparative Analysis of Baseline and Optimized Models*

The optimized SVM model outperformed the base model, particularly for minority intents. Figure ?? displays the confusion matrix for the optimized model, highlighting the improvements in precision and recall for challenging classifications.

4.7 Conclusion

Random Forest and SVM both demonstrated strong performance in this experiment. Random Forest's ensemble nature provided robustness and generalization, making it the best-performing model overall. SVM, particularly with the linear kernel, offered comparable results but required more tuning to handle imbalanced classes effectively. These findings highlight the importance of selecting models and hyperparameters tailored to the dataset's characteristics and the problem's requirements.

4.8 Feature Experimentation

...




Fig. 8. Random Forest Confusion Matrices




Fig. 9. Random Forest Evaluation Metrics




Fig. 10. Random Forest Weighted F1 Scores




Fig. 11. Random Forest Classification Report Train Set




Fig. 12. Random Forest Classification Report Test Set

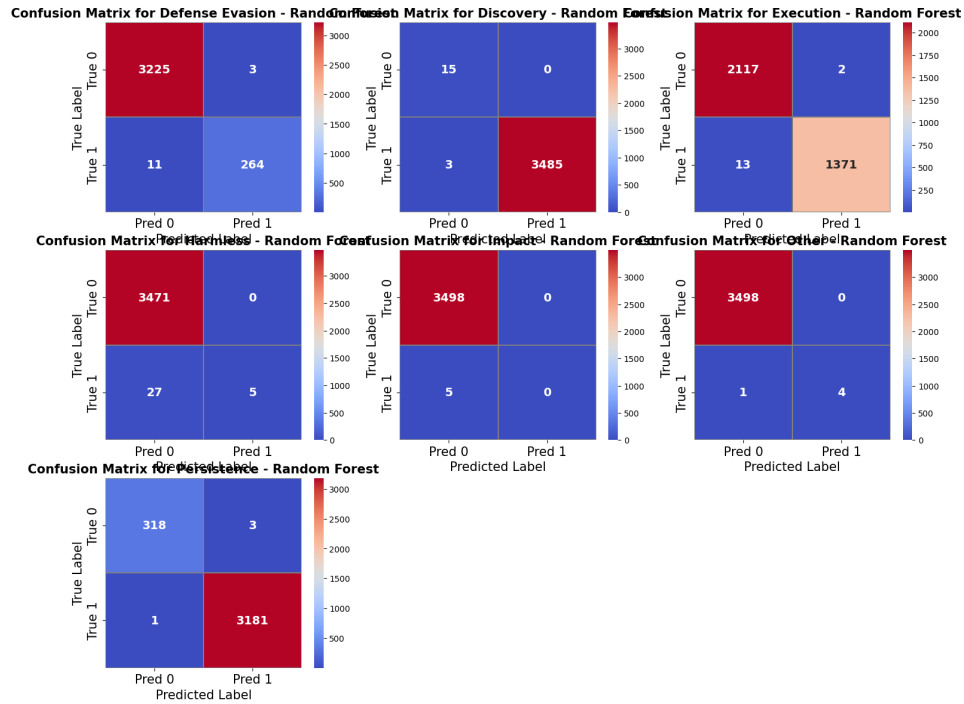


Fig. 13. SVM Confusion Matrices




Fig. 14. SVM Evaluation Metrics





Fig. 15. SVM Weighted F1 Scores



Publication date: January 2025.

Fig. 16. SVM Classification Report Train Set




Fig. 17. SVM Classification Report Test Set

5 UNSUPERVISED LEARNING - CLUSTERING

5.1 Introduction

This section of the project focuses on applying unsupervised learning techniques to analyze SSH attack data. Specifically, clustering methods were utilized to uncover patterns and relationships within the dataset. An emphasis was placed on identifying distinct attack behaviors and intents.

5.2 Data Preparation

The dataset chosen was the one generated through the TF-IDF vectorization technique. This was made because it was essential to start with a dataset that represented in the best way the frequency and the importance of words, making each word as a dimension of our vector.

5.3 Clustering Methods

Two clustering techniques were applied:

- **K-Means Clustering:** Using the Elbow Method and silhouette scores, the optimal number of clusters was determined.
- **Gaussian Mixture Model (GMM):** Similar to K-Means, the optimal number of clusters was selected using silhouette scores and log-likelihood values.

5.4 Hyperparameter Tuning

Hyperparameter tuning was conducted to optimize the performance of both clustering methods. For K-Means, parameters such as the initialization method (k-means++ and random), the number of initializations, and the maximum number of iterations were fine-tuned using a grid search approach. Similarly, GMM parameters including the initialization method (kmeans), covariance type (full and spherical), and tolerance were optimized. These steps ensured the models were tailored to the dataset, resulting in better clustering outcomes.

5.5 Clusters Visualization

To visualize the clustering results, t-SNE dimensionality reduction was applied. Its functionalities makes it an excellent choice for visualizing clusters in datasets where direct interpretation is difficult due to high dimensionality. By projecting the data into a two-dimensional space, t-SNE enables us to identify patterns and groupings that may not be evident in the original feature space. The two-dimensional plots provided a clear representation of the clusters formed by both K-Means and GMM, highlighting their separability and internal consistency.

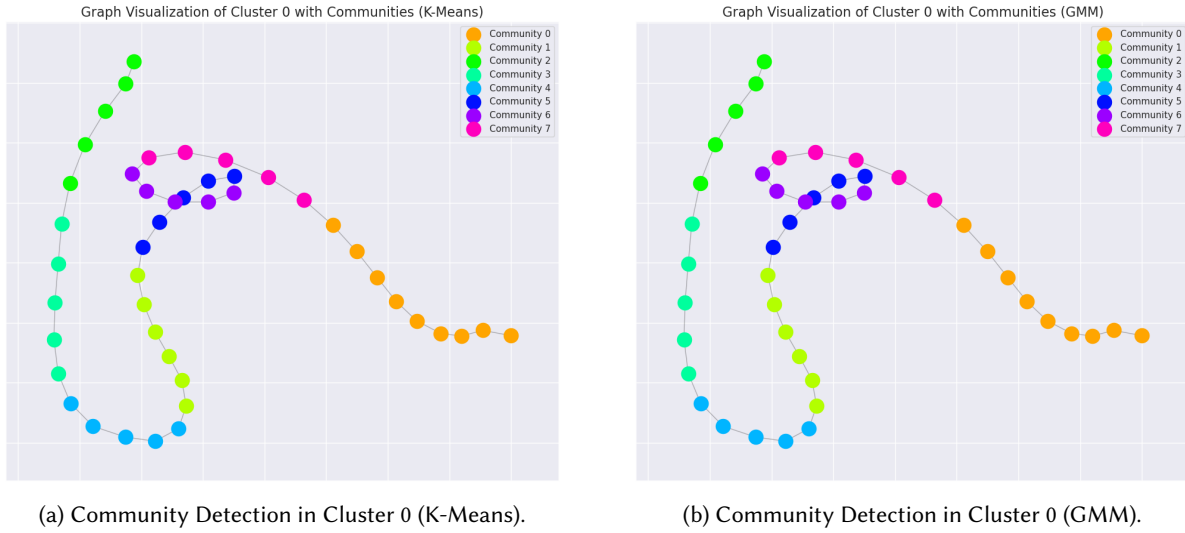


Fig. 20

5.7 Conclusion

This analysis provided valuable insights into the patterns of SSH attacks through clustering. The K-Means and GMM algorithms both effectively identified meaningful clusters, as supported by validation metrics and visualizations. Hyperparameter tuning further enhanced the performance of both models. The results demonstrate the potential of unsupervised learning in uncovering hidden patterns in complex datasets, providing a foundation for future applications such as anomaly detection and improved cybersecurity strategies.

6 LANGUAGE MODEL EXPLORATION

6.1 Introduction

This report documents the steps taken to fine-tune a pre-trained BERT model for multi-label classification using a dataset of SSH attack sessions. The focus was on customizing the final classification layer and training the model on the 'Set_Fingerprint' intents. Performance metrics and visualizations are presented to evaluate the effectiveness of the model.

6.2 Data Preprocessing

The dataset was processed as follows:

- The 'Set_Fingerprint' column was preprocessed for multi-label encoding using the 'MultiLabelBinarizer' class, enabling the representation of intents as binary vectors.
- The dataset was split into 70% training, 20% validation, and 10% testing sets.

6.3 Tokenization

Pre-trained BERT's tokenizer was used to convert text data into input IDs and attention masks. Tokenization results were saved to disk for reuse, reducing computational overhead.

6.4 Model Architecture

A pre-trained ‘bert-base-uncased’ model was used, with a custom linear classification head added to predict the intents. The classifier had an input size matching BERT’s hidden layer dimensions and an output size equal to the number of intents.

6.5 Training Process

- The model was trained using the AdamW optimizer with a learning rate of 5×10^{-5} and a linear scheduler.
- A ‘BCEWithLogitsLoss’ loss function was applied for multi-label classification.
- Mixed precision training was used to enhance computational efficiency.
- Training was conducted for 4 epochs, using a gradient accumulation strategy to manage memory constraints.

6.6 Evaluation Metrics

- Metrics included precision, recall, F1-score, and ROC-AUC.
- ROC curves and training/validation loss curves were generated for each class.

6.7 Results

6.7.1 Training and Validation Loss

The training loss decreased steadily, indicating effective learning. However, the validation loss started increasing after epoch 2, suggesting potential overfitting. (See Figure 21).




Fig. 21. Training and Validation Loss

6.7.2 ROC Curves

ROC curves were generated for each intent class. The area under the curve (AUC) values were as follows:

- **High AUC:** Discovery (1.00), Execution (0.98), Defense Evasion (0.97), Persistence (0.94).
- **Low AUC:** Harmless (0.77), Other (0.48), Impact (NaN due to lack of positive samples).

The variability in AUC values indicates that the model performs well on some intents but struggles with others. (See Figure 22).




Fig. 22. ROC Curves for Each Class

6.7.3 Test Metrics

The model achieved the following metrics on the test set:

- Precision: 0.94
- Recall: 0.92
- F1-Score: 0.93
- Weighted ROC-AUC: 0.91

6.8 Discussion

The results demonstrate that the fine-tuned BERT model performs well for most classes, with high AUC scores for key intents like 'Discovery' and 'Execution'. However, challenges remain:

- The model struggled with the 'Other' class, likely due to limited or ambiguous data.
- Overfitting was observed after epoch 2, necessitating regularization techniques like dropout or early stopping in future iterations.

6.9 Conclusion

Fine-tuning a pre-trained BERT model proved effective for multi-label classification of SSH attack intents. While the model achieved high overall performance, further tuning and data augmentation could address issues with minority classes and overfitting.

7 CONCLUSION

7.1 Summary of Key Findings

In this project, we explored various techniques for analyzing and classifying SSH shell attack logs. The primary objectives were to preprocess the data, perform exploratory data analysis, implement supervised and unsupervised learning models, and leverage advanced language models for classification tasks. Here, we summarize the key findings from each section of the project.

Data Exploration and Pre-processing: We began by loading and inspecting the dataset, identifying missing values, and handling duplicates. Temporal analysis revealed significant variations in attack frequencies over time, with notable peaks during specific hours and months. Feature extraction and common words analysis provided insights into the most frequent commands and intents used in the attack sessions.

Supervised Learning - Classification: We implemented and evaluated several machine learning models, including Logistic Regression, Random Forest, and Support Vector Machine (SVM). Hyperparameter tuning improved the performance of these models, and the result analysis highlighted the strengths and weaknesses of each approach. Feature experimentation with different text representation techniques, such as Bag of Words (BoW) and TF-IDF, demonstrated the impact of feature selection on model performance.

Unsupervised Learning - Clustering: Clustering techniques, such as K-Means and Gaussian Mixture Models (GMM), were used to group similar attack sessions. The elbow method and silhouette analysis helped determine the optimal number of clusters. Cluster visualization using t-SNE provided a clear representation of the clusters, and cluster analysis revealed common patterns and behaviors within each group.

Language Model Exploration: We explored the use of advanced language models, such as BERT, for classifying attack session tactics. Fine-tuning the pretrained BERT model on our dataset improved classification performance. Learning curves indicated the optimal number of epochs for training, helping to avoid overfitting.

7.2 Challenges Faced

Throughout the project, we encountered several challenges that required careful consideration and problem-solving.

Data Quality and Preprocessing: Handling missing values, duplicates, and inconsistencies in the dataset was a critical step. Ensuring the data was clean and well-prepared for analysis required significant effort. Additionally, the unstructured nature of the session text posed challenges for text representation and feature extraction.

Model Selection and Tuning: Selecting appropriate machine learning models and tuning their hyperparameters was a complex task. Balancing model complexity with performance and avoiding overfitting required iterative experimentation and validation.

Computational Resources: Training advanced language models, such as BERT, required substantial computational resources. Efficiently managing these resources and optimizing the training process was essential to achieve timely results.

Interpretability of Results: Interpreting the results of clustering and classification models, especially in the context of cybersecurity, was challenging. Ensuring that the findings were meaningful and actionable required careful analysis and domain knowledge.

7.3 Future Work

Based on the findings and challenges encountered in this project, we propose several directions for future work.

Enhanced Feature Engineering: Further exploration of feature engineering techniques, such as incorporating domain-specific knowledge and using advanced text representation methods, could improve model performance. Experimenting with additional features, such as network metadata and contextual information, may provide deeper insights into attack patterns.

Advanced Model Architectures: Exploring more advanced model architectures, such as transformer-based models and deep neural networks, could enhance classification accuracy. Transfer learning with other pretrained models and ensemble methods may also yield better results.

Real-time Analysis and Detection: Implementing real-time analysis and detection systems for SSH shell attacks could provide immediate insights and responses to potential threats. Integrating the models developed in this project into a real-time monitoring framework would be a valuable extension.

Broader Dataset and Generalization: Expanding the dataset to include a wider range of attack types and sources would improve the generalizability of the models. Collaborating with other organizations to share data and insights could enhance the robustness and applicability of the findings.

7.4 Conclusion

This project demonstrated the potential of machine learning and advanced language models for analyzing and classifying SSH shell attack logs. By leveraging various techniques, we gained valuable insights into attack patterns and behaviors, which can inform cybersecurity strategies and defenses. Despite the challenges faced, the results highlight the importance of data-driven approaches in enhancing cybersecurity threat detection and response capabilities. Future work in this area holds promise for further advancements and practical applications in the field of cybersecurity.