

SSH Shell Attacks

ANDREA BOTTICELLA*, Politecnico di Torino, Italy

ELIA INNOCENTI*, Politecnico di Torino, Italy

RENATO MIGNONE*, Politecnico di Torino, Italy

SIMONE ROMANO*, Politecnico di Torino, Italy

This paper introduces a comprehensive machine learning approach to analyze SSH shell attack sessions, leveraging both supervised and unsupervised learning techniques. Using a dataset of 230,000 unique Unix shell attack sessions, the methodology aims to classify attacker tactics based on the MITRE ATT&CK framework and uncover latent patterns through clustering. The key contributions of this work are:

- Development of a robust pre-processing pipeline to analyze temporal trends, extract numerical features, and evaluate intent distributions from large-scale SSH attack session data.
- Implementation of supervised classification models to accurately predict multiple attacker tactics, supported by hyperparameter tuning and feature engineering for enhanced performance.
- Application of unsupervised clustering techniques to uncover hidden patterns in attack behaviors, leveraging visualization tools and cluster analysis for fine-grained categorization.
- Exploration of advanced language models, such as BERT or Doc2Vec, for representation learning and fine-tuning to improve intent classification and session interpretation.

CCS Concepts: • **Computing methodologies** → Supervised learning by classification; Unsupervised learning; Natural language processing; Machine learning; Machine learning approaches; • **Security and privacy** → Intrusion detection systems.

Additional Key Words and Phrases: Machine learning, supervised learning, unsupervised learning, language models, text classification, clustering, intent classification, SSH shell attacks, security log analysis

CONTENTS

Abstract	1
Contents	1
1 INTRODUCTION	1
2 BACKGROUND	2
3 DATA EXPLORATION AND PRE-PROCESSING	3
4 SUPERVISED LEARNING - CLASSIFICATION	8
5 UNSUPERVISED LEARNING - CLUSTERING	14
6 LANGUAGE MODEL EXPLORATION	18
7 CONCLUSION	21

1 INTRODUCTION

1.1 Motivation

Security logs play a crucial role in understanding and mitigating cyber attacks, particularly in the domain of network and system security. With the increasing sophistication of cyber threats, analyzing and interpreting security logs has become paramount for detecting, preventing, and responding to potential security breaches.

*The authors collaborated closely in developing this project.

Authors' Contact Information: Andrea Botticella, andrea.botticella@studenti.polito.it, Politecnico di Torino, Turin, Italy; Elia Innocenti, elia.innocenti@studenti.polito.it, Politecnico di Torino, Turin, Italy; Renato Mignone, renato.mignone@studenti.polito.it, Politecnico di Torino, Turin, Italy; Simone Romano, simone.romano2@studenti.polito.it, Politecnico di Torino, Turin, Italy.

Unix shell attacks, especially those executed through SSH protocol, represent a significant vector for potential system compromises.

The complexity of security log analysis stems from several key challenges:

- Logs are often unstructured and contain ambiguous or malformed text.
- Manual parsing and interpretation of logs is time-consuming and error-prone.
- The sheer volume of log data makes comprehensive manual review impractical.

These challenges underscore the need for automated, intelligent approaches to log analysis that can efficiently extract meaningful insights and identify potential security threats.

1.2 Objective

The primary objective of this research is to develop and evaluate machine learning techniques for automatic analysis and classification of SSH shell attack logs. Specifically, we aim to:

- Automate the process of log analysis and intent classification.
- Provide security professionals with insights into attack strategies.
- Identify patterns and trends in attack sessions for proactive threat mitigation.

The significance of this research lies in its potential to enhance cybersecurity threat detection and response capabilities by transforming complex, unstructured log data into actionable intelligence.

2 BACKGROUND

2.1 Security Logs and Attack Analysis

In the context of SSH shell attacks, logs document the sequence of commands executed during a malicious session, enabling security researchers to analyze attacker behaviors, techniques, and potential system impacts. However, due to the high quantity of logs generated by security systems, manual analysis is impractical, necessitating automated approaches and valid frameworks to extract meaningful insights.

2.2 MITRE ATT&CK Framework

The MITRE ATT&CK (Adversarial Tactics, Techniques, and Common Knowledge) framework provides a comprehensive knowledge base of adversary tactics and techniques observed in real-world cyber attacks. This framework serves as a standardized methodology for understanding and categorizing attack strategies.

For our research, we focus on seven key intents derived from the MITRE ATT&CK framework:

- **Persistence:** Techniques used by adversaries to maintain system access across restarts or credential changes.
- **Discovery:** Methods for gathering information about the target system and network environment.
- **Defense Evasion:** Strategies to avoid detection by security mechanisms.
- **Execution:** Techniques for running malicious code on target systems.
- **Impact:** Actions aimed at manipulating, interrupting, or destroying systems and data.
- **Other:** Less common tactics including Reconnaissance, Resource Development, Initial Access, etc.
- **Harmless:** Non-malicious code or actions.

2.3 Research Approach

Our research employs a multi-faceted approach to this analysis:

- Explore and preprocess a large dataset of Unix shell attack sessions.
- Apply supervised learning techniques to classify attack tactics based on session characteristics.
- Utilize unsupervised learning methods to discover patterns and clusters in attack sessions.

- Investigate the potential of advanced language models in understanding and categorizing attack intents.

3 DATA EXPLORATION AND PRE-PROCESSING

3.1 Introduction

This section outlines the steps taken to explore and preprocess the dataset used in this study. The primary objective is to understand the data characteristics, identify patterns, and prepare the data for further analysis. We focus on temporal trends, intent distributions, and textual features.

3.2 Dataset Overview

The dataset consists of logs from SSH attacks. Each entry includes the following features:

- **session_id**: An integer representing the unique identifier for each session.
- **full_session**: A string containing the full sequence of commands executed during the attack session.
- **first_timestamp**: The timestamp indicating the start of the session.
- **Set_Fingerprint**: A set of strings (labels) describing the nature of the attack.

Below is an example of the dataset structure:

	session_id	full_session	first_timestamp	Set_Fingerprint
0	0	enable; system; shell; sh; cat /proc/mounts; /...	2019-06-04 09:45:11.151186+00:00	[Defense Evasion, Discovery]
1	1	enable; system; shell; sh; cat /proc/mounts; /...	2019-06-04 09:45:50.396610+00:00	[Defense Evasion, Discovery]
2	2	enable; system; shell; sh; cat /proc/mounts; /...	2019-06-04 09:54:41.863315+00:00	[Defense Evasion, Discovery]
...
233034	233046	cat /proc/cpuinfo grep name wc -l; echo -e...	2020-02-29 23:59:22.199490+00:00	[Discovery, Persistence]

233035 rows × 4 columns

Fig. 1. Dataset

3.3 Dataset Preparation and Feature Extraction

The dataset used in this study was provided as a Parquet file.

To prepare the dataset for analysis and ensure its quality, several preprocessing steps were undertaken, described in detail below:

- (1) **Decoding the full_session column**: The `full_session` column contained 90026 encoded shell scripts, making the raw data difficult to interpret. A decoding process was applied to these entries, converting them into a human-readable format.
- (2) **Formatting the first_timestamp column**: The `first_timestamp` column was checked for consistency and converted into a standard datetime format.
- (3) **Splitting the full_session column into lists of commands**: The `full_session` column contained entire attack sessions represented as single strings. Since we wanted each word / command to be a feature of our models, it was necessary to break it down into individual commands or keywords for more granular analysis. This was achieved through a multi-step splitting process using specific delimiters:
 - The semicolon (;) was used to separate distinct commands within a session.
 - The pipe (|) was used to divide concatenated commands or pipelines.
 - Whitespace () was used to further split commands into individual words or arguments.

This process allowed the identification of specific actions and parameters, facilitating a detailed analysis of the attack strategies.

- (4) **Cleaning individual commands:** Each command or keyword extracted from the `full_session` column was processed to remove unnecessary or undesired elements. Regular expressions were used to strip unwanted symbols, and variable assignments.
- (5) **Handling missing values:** To ensure the integrity and completeness of the dataset, a thorough check for missing values was performed across all columns.

```
enable ; system ; shell ; sh ; cat /proc/mounts ; /bin/busybox SAEMW ; cd /dev/
shm ; cat .s || cp /bin/echo .s ; /bin/busybox SAEMW ; tftp ; wget ; /bin/
busybox SAEMW ; dd bs=52 count=1 if=.s || cat .s || while read i ; do echo
$i ; done < .s ; /bin/busybox SAEMW ; rm .s ; exit ;

[enable, system, shell, sh, cat, /proc/mounts, /bin/busybox, SAEMW, cd, /dev/shm
, cat, .s, cp, /bin/echo, .s, /bin/busybox, SAEMW, tftp, wget, /bin/busybox
, SAEMW, dd, bs, count, if, cat, .s, while, read, i, do, echo, i, done, .s,
/bin/busybox, SAEMW, rm, .s, exit]
```

Listing 1. Dataset Processing

These preprocessing steps ensured that the dataset was well-structured, clean, standardized and ready for advanced analysis and machine learning tasks. By addressing issues such as encoding, formatting, and missing data, the preprocessing phase established a robust foundation for the study.

3.4 Temporal Analysis

The temporal analysis examines when the attacks were performed and the intent distribution over time.

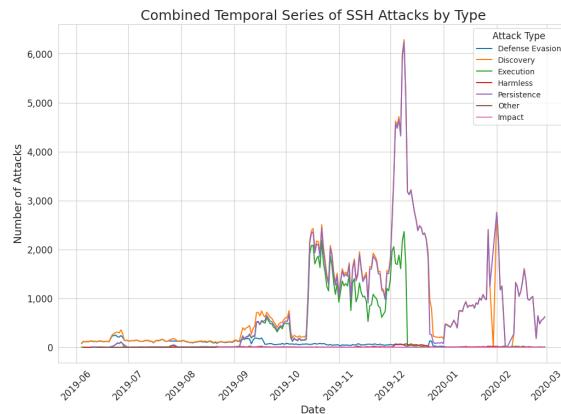


Fig. 2. Temporal Series of SSH Attacks

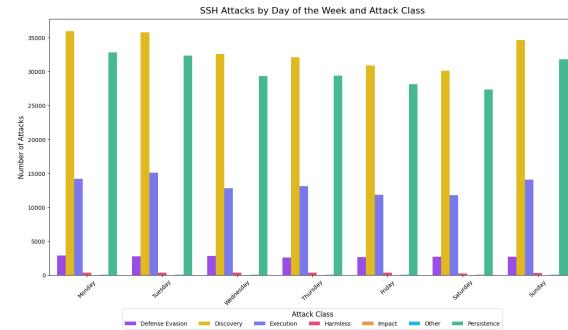


Fig. 3. Attacks by the day of the week

Our dataset spans from June 2019 to March 2020, revealing a notable concentration of attacks between mid-October 2019 and January 2020.

Further analysis was conducted on the distribution of sessions across various temporal dimensions (months, days of the week, etc.), but no significant patterns were identified.

3.5 Common Words Analysis

The analysis of common words within the attack sessions provides insights into frequently utilized commands and keywords. These elements are pivotal for understanding attacker behavior and informing model training for detecting attack intents.

To visualize and analyze the most frequent words, two approaches were used:

- **Word Cloud Visualization:** A word cloud (Figure 4) was generated to represent the most common words in the dataset. Commands like `grep`, `cat`, `echo`, and `rm` dominate the word cloud, highlighting their frequent usage in attack sessions.
- **Top 20 Most Common Words Bar Plot:** A bar plot (Figure 5) showcases the top 20 most frequent words, ordered by their frequency. This plot provides more quantitative detail, revealing the high occurrence of commands such as `grep` (over 1.2 million occurrences) and `cat` (over 1 million occurrences). Other significant commands include `echo`, `rm`, `uname`, and `name`. These commands are commonly associated with file manipulation, string searching, and process information, which are typical in malicious activities.

The combination of these visualizations supports the identification of commonly used commands and their potential roles in attack sessions.

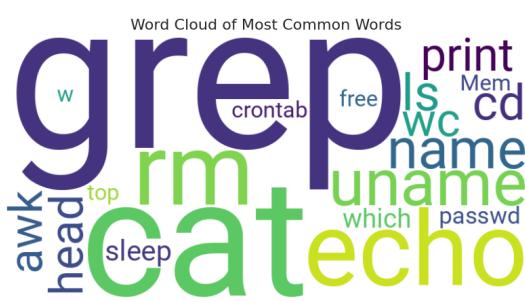


Fig. 4. Word Cloud of Most Common Words

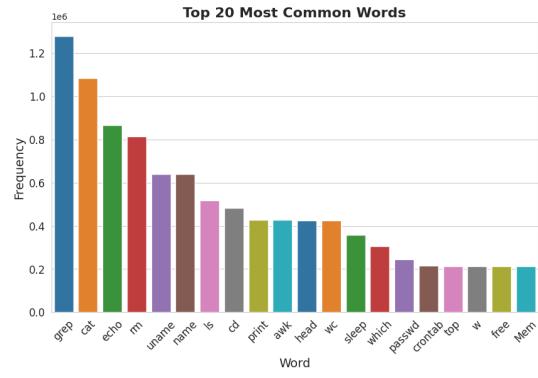


Fig. 5. Top 20 Most Common Words

3.6 Intent Analysis

This section presents an analysis of different intents identified in attack sessions and their relationships. Our analysis reveals distinct patterns in both the frequency of individual intents and their co-occurrence relationships, providing valuable insights into attacker behavior patterns.

Figure 6 illustrates the frequency distribution of different intents in our dataset. The analysis reveals that Discovery and Persistence are the most prevalent intents, with approximately 200,000 occurrences each. This suggests that attackers primarily focus on reconnaissance activities and establishing long-term presence in the targeted systems. Execution intent appears as the third most common, with roughly 100,000 occurrences, indicating a significant number of attempts to run malicious code. Defense Evasion shows notably fewer occurrences (around 20,000), while Harmless, Other, and Impact intents are relatively rare in the dataset.

The co-occurrence heatmap (Figure 7) reveals significant patterns in how different intents interact within attack sessions. Several key observations emerge:

- The strongest co-occurrence relationship exists between Discovery and Persistence (211,281 co-occurrences), suggesting that attackers often combine reconnaissance activities with attempts to maintain system access.

- Execution shows strong correlations with both Discovery (92,279) and Persistence (91,576), indicating that malicious code execution frequently accompanies both system discovery and persistence establishment attempts.
- Defense Evasion exhibits moderate co-occurrence with Discovery (18,825) and minimal correlation with other intents, suggesting that evasion techniques are primarily employed during reconnaissance phases.
- Harmless, Impact, and Other intents show minimal co-occurrence with other categories (with values mostly under 20 co-occurrences). This isolation pattern is partially explained by their low frequency in the dataset: Harmless with 2,206 occurrences, Impact with only 27 occurrences, and Other with 327 occurrences. These numbers represent a small fraction of the total recorded intents, naturally limiting their potential for co-occurrence with other intent types.

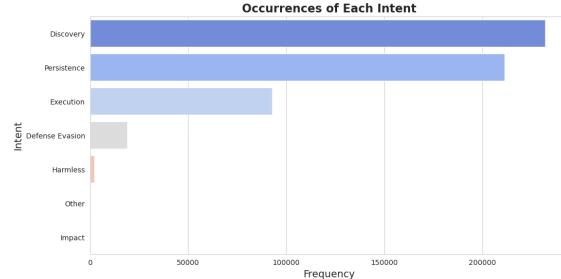


Fig. 6. Occurrences of each Intent

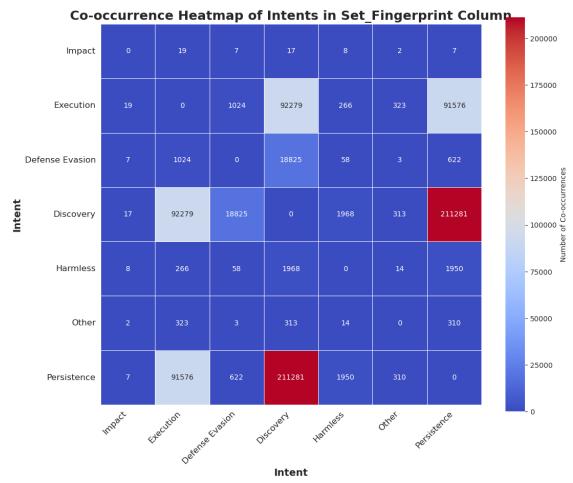


Fig. 7. Co-Occurrence Heatmap of Intents

These findings highlight common attack patterns where adversaries typically begin with discovery operations, followed by persistence establishment and execution of malicious code. This understanding can inform the development of detection strategies and defensive measures, particularly focusing on the most common intent combinations.

3.7 Session Analysis

The session analysis aims to understand the structural characteristics of the attack sessions in the dataset. This is accomplished through two visualizations: the Empirical Cumulative Distribution Function (ECDF) plots and the distribution of the number of words per session. The ECDF plots in Figure 8 provide insights into the distribution of the number of characters and words across all sessions.

- The ECDF for characters indicates that the majority of sessions have fewer than 10^3 characters, with a sharp increase between 10^2 and 10^3 characters. This suggests a high concentration of relatively short sessions.
- The ECDF for words shows a similar trend, with most sessions containing fewer than 10^2 words. The distribution reflects the concise nature of many attack sessions, potentially focusing on executing a limited number of commands.

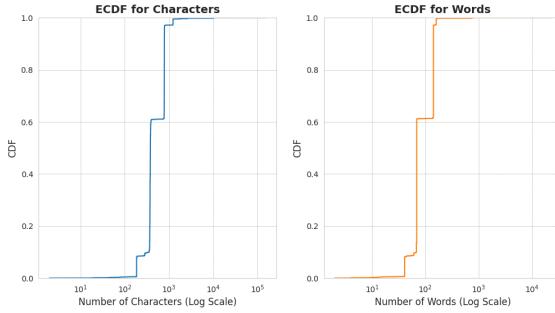


Fig. 8. ECDF for Characters and Words in Sessions

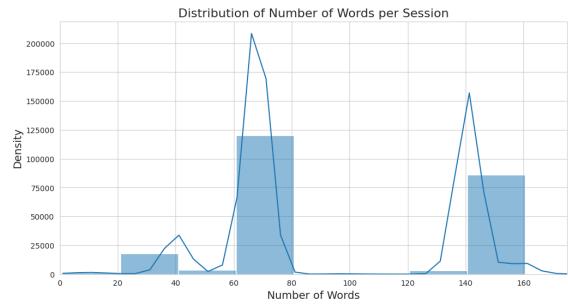


Fig. 9. Distribution of the Number of Words per Session

Figure 9 illustrates the density distribution of the number of words per session. Key observations include:

- The distribution is bimodal, with two distinct peaks. This likely reflects different categories of sessions, such as exploratory attacks with a larger number of commands versus simpler, targeted attacks with fewer commands.
- The majority of sessions have fewer than 100 words, reinforcing the compactness observed in the ECDF plots.
- Outliers with higher word counts likely represent complex sessions, potentially involving multiple stages or more sophisticated attack strategies.

The relevance of these plots lies in their ability to provide a foundational understanding of session structure, which is critical for feature engineering and model development. Shorter sessions may correspond to quick, automated attacks, while longer sessions might represent manual or multi-step intrusions. These insights can inform the design of models by emphasizing features tailored to the varying lengths and complexities of sessions, ultimately improving detection accuracy.

3.8 Text Representation

To enable further analysis of the dataset, the initial textual data was transformed into numerical representations. This step is essential for applying machine learning techniques, as numerical formats are required for model training and evaluation. Two widely used text representation techniques were employed:

- **Bag of Words (BoW):** This technique represents each session as a vector where each dimension corresponds to a specific word, and the value represents the frequency of that word in the session. While simple and interpretable, BoW does not capture the importance or uniqueness of words across the dataset.
- **Term Frequency-Inverse Document Frequency (TF-IDF):** This method extends BoW by weighting the word frequencies based on their inverse document frequency, thereby emphasizing words that are important within a session but less frequent across all sessions. This approach helps us, through the "min_df" parameter to remove from the initial dataset, the words that have a very low frequency. In our case, we set the "min_df" parameter to 0.01, which means that words that appear in less than 1% of the sessions are removed. In this way the features are reduced from 300,000 to 90.

Figure 10 and Figure 11 illustrate the transformed datasets using the BoW and TF-IDF techniques, respectively. The TF-IDF representation was chosen for further analysis as it captures meaningful patterns by emphasizing word relevance, ensuring that subsequent analyses focus on the most distinctive features of each session, thereby enhancing the effectiveness of machine learning models in identifying patterns and predicting intents.

	Set_Fingerprint	0kx34uax1rv	15s	172	192	20m	3s	8m	9p7vd0epz3tz	admin	...	up	var	wc	wget	which	while	x13	x17	x19	xf
0	[Defense Evasion, Discovery]	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	...	0.000000	0.000000	0.000000	0.000559	0.000000	0.007507	0.0	0.0	0.0	0.0
1	[Defense Evasion, Discovery]	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	...	0.000000	0.000000	0.000000	0.000559	0.000000	0.007507	0.0	0.0	0.0	0.0
2	[Defense Evasion, Discovery]	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	...	0.000000	0.000000	0.000000	0.000559	0.000000	0.007507	0.0	0.0	0.0	0.0
...
233034	[Discovery, Persistence]	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.002631	...	0.001042	0.001611	0.002168	0.000000	0.001432	0.000000	0.0	0.0	0.0	0.0

233035 rows × 88 columns

Fig. 10. Bag-of-Words

	Set_Fingerprint	0kx34uax1rv	15s	172	192	20m	3s	8m	9p7vd0epz3tz	admin	...	up	var	wc	wget	which	while	x13	x17	x19	xf
0	[Defense Evasion, Discovery]	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	...	0.000000	0.000000	0.000000	0.154131	0.000000	0.156517	0.0	0.0	0.0	0.0
1	[Defense Evasion, Discovery]	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	...	0.000000	0.000000	0.000000	0.154131	0.000000	0.156517	0.0	0.0	0.0	0.0
2	[Defense Evasion, Discovery]	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	...	0.000000	0.000000	0.000000	0.154131	0.000000	0.156517	0.0	0.0	0.0	0.0
...
233034	[Discovery, Persistence]	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.010208	...	0.067277	0.465673	0.133076	0.000000	0.066583	0.000000	0.0	0.0	0.0	0.0

233035 rows × 88 columns

Fig. 11. TF-IDF

4 SUPERVISED LEARNING - CLASSIFICATION

4.1 Introduction

The supervised learning experiment in this project aimed to classify attack sessions into various intent categories derived from the `Set_Fingerprint` column of the dataset. This section explores the use of Random Forest and Support Vector Machines (SVM) as the primary models. The analysis focuses on how these models handle multi-label classification and evaluates their performance using metrics such as weighted F1-scores and confusion matrices.

This section outlines the model training and evaluation processes, and a detailed discussion of the results obtained. Each model's strengths and weaknesses are analyzed, providing insights into their application to multi-label classification problems.

4.2 Data Preprocessing

To effectively apply supervised learning models, it is crucial to represent the textual data in a numerical format. Raw text cannot be directly processed by most machine learning algorithms, so as we said in the previous section, we transformed the dataset into a structured numerical form.

Unlike BoW, which merely counts word occurrences without differentiating their relevance, TF-IDF down-weights frequently occurring words that may not carry significant information (e.g., common command-line syntax) and up-weights rare but potentially more meaningful terms. This property helps improve the model's ability to differentiate between different types of attacks and intents.

To prepare the data for supervised learning:

- (1) **Encoding Intents:** After loading the TF-IDF dataset, the `Set_Fingerprint` column was encoded into multi-label binary format using the MultiLabelBinarizer. Each intent was represented as a binary vector, allowing for simultaneous prediction of multiple labels.
- (2) **Splitting the Data:** The dataset was divided into training (70%), validation (20%), and testing (10%) subsets. Stratified splitting was employed to ensure that all intents were well-represented in each subset.

These steps ensured the dataset was clean, balanced, and ready for supervised learning.

4.3 Model Training

Three models were trained and evaluated using their default configurations to establish baseline performance:

1. Random Forest

The Random Forest model was trained with default parameters, including 100 estimators and unlimited maximum depth. This initial training provided insights into potential overfitting or underfitting issues and served as a benchmark for subsequent tuning.

2. Support Vector Machines (SVM)

SVM was initially trained with default settings using a linear kernel and a regularization parameter $C = 1$. The performance was evaluated to assess the model's ability to handle multi-label classification tasks with linearly separable data.

3. Logistic Regression

In order to have another view of the analysis, we performed the training with the logistic regression model. While the model performed well, the results were not as significant as those of RF and SVM, so they will not be discussed in detail in this section (see Appendix).

4.4 Evaluation Metrics

The models were evaluated using the following metrics:

- **Accuracy, Precision, Recall:** Basic evaluation metrics.
- **Confusion Matrices:** Provided insight into TP, FP, FN, and TN for each intent.
- **Weighted F1-Scores:** Measured the harmonic mean of precision and recall, with weights proportional to class support.

This evaluation allowed for the identification of baseline performance, highlighting potential areas for improvement through hyperparameter tuning.

4.5 Hyperparameter Tuning

To optimize each model, hyperparameter tuning was performed using a grid search approach. This process aimed to improve performance and address issues of overfitting or underfitting observed in the baseline models:

1. Random Forest

The grid search explored combinations of the number of estimators (50, 100, and 150) and maximum depth (10, 50, and 100). The best-performing configuration was selected based on weighted F1-scores.

2. Support Vector Machines (SVM)

For SVM, the grid search varied the regularization parameter C (0.1, 1, 10, 100) and the kernel type (linear and RBF). Additional tuning for the RBF kernel included the gamma parameter (scale and auto).

4.6 Results and Observations

4.6.1 Random Forest

• Performance Overview with base model

The Random Forest model demonstrated exceptional performance across most attack classifications. On the test set, it achieved remarkable weighted average scores with precision of 0.999, recall of 0.994, and F1-score of 0.996. The model showed particular strength in classifying Discovery attacks, achieving perfect scores (1.000) across precision, recall, and F1-score, with substantial support of 69,659 samples.

Notable performance metrics for major attack categories include:

- Defense Evasion: precision (0.993), recall (0.969), F1-score (0.981)
- Execution: precision (0.999), recall (0.988), F1-score (0.994)
- Persistence: precision (0.999), recall (1.000), F1-score (1.000)

However, the model showed some limitations with the Harmless and Impact categories, achieving lower performance metrics:

- Harmless: precision (0.939), recall (0.160), F1-score (0.273)
- Impact: precision (0.500), recall (0.500), F1-score (0.500)

The model's strength lies in its ability to maintain high precision and recall across most attack categories, particularly for well-represented classes. However, the confusion matrices reveal an interesting pattern: the model shows some limitations in classifying Harmless and Impact categories, indicating potential challenges in distinguishing benign command sequences from malicious ones.

- *Hyperparameter Analysis*

Hyperparameter tuning improved the Random Forest model's performance, particularly for minority intents. Figure 14 highlights the improvement in weighted F1-scores across different configurations.

The confusion matrices reveal the model's classification behavior in detail. For Defense Evasion, out of 69,911 total cases:

- True Negatives: 64,241 cases | True Positives: 5,460 cases
- False Positives: 174 cases | False Negatives: 36 cases

The model showed particularly strong performance in Discovery classification, with only 19 total misclassifications (10 false positives and 9 false negatives) out of 69,911 cases, demonstrating exceptional precision in identifying this attack type. The optimized model showed improved handling of edge cases while maintaining strong performance on mainstream attack patterns. A particularly noteworthy observation is how the model balances precision and recall across different attack categories. The tuning process helped achieve a better equilibrium, especially for categories with smaller representation in the dataset.

- *Comparative Analysis of Baseline and Optimized Models*

Figure 15 and 16 present the confusion matrix for the optimized Random Forest model. Comparing the base and optimized Random Forest models reveals a fascinating pattern of improvements. The optimized model shows more nuanced decision boundaries, particularly evident in the handling of edge cases. As shown in the performance metrics, the gap between training and testing performance narrowed, indicating better generalization capabilities.

The model's robustness is particularly evident in its consistent performance across both training and test sets, suggesting effective learning of underlying patterns rather than mere memorization. This consistency is crucial in security applications where new, slightly varied attack patterns must be detected reliably.

4.6.2 Support Vector Machines (SVM)

- *Performance Overview with base model*

The SVM base model performed well for linearly separable data, achieving high precision and recall for intents with large sample sizes. However, it struggled with minority intents due to its sensitivity to class imbalances. Figure 17 illustrates the confusion matrix for the base model.

The weighted averages on the test set showed excellent overall performance:

- Precision: 0.999 | Recall: 0.994 | F1-score: 0.995

Particularly strong performance was observed in:

- Defense Evasion: precision (0.993), recall (0.983), F1-score (0.988)
- Execution: precision (0.996), recall (0.994), F1-score (0.995)

- Persistence: precision (0.998), recall (0.996), F1-score (0.997)

As evidenced in Figure 17, the model's strength lies in its ability to establish clear decision boundaries for well-defined attack categories. However, the performance matrices reveal an interesting limitation: the model shows more pronounced difficulties with the Impact category compared to Random Forest, suggesting that linear separation might not be optimal for all attack types.

- *Hyperparameter Analysis*

Grid search tuning improved SVM's performance, especially for the RBF kernel. Figure 19 shows the weighted F1-scores for different hyperparameter combinations. The confusion matrices for SVM reveal slightly different classification patterns compared to Random Forest:

For Defense Evasion:

- True Negatives: 64,220 (slightly lower than RF) | True Positives: 5,445 (slightly lower than RF)
- False Positives: 189 (higher than RF's 174) | False Negatives: 57 (higher than RF's 36)

A notable difference appeared in Impact classification, where SVM showed significantly lower performance:

- Precision: 0.500 | Recall: 0.500 | F1-score: 0.500

The confusion matrices post-tuning reveal an interesting shift in classification patterns. While the overall accuracy remained high, the distribution of errors changed, suggesting that the optimized model developed different decision boundaries that better reflect the natural grouping of attack patterns in the feature space.

- *Comparative Analysis of Baseline and Optimized Models*

The optimized SVM model outperformed the base model, particularly for minority intents. Figure 20 and 21 display the confusion matrix for the optimized model, highlighting the improvements in precision and recall for challenging classifications.

The evolution from baseline to optimized SVM model reveals important insights about the nature of attack classification. The optimized model shows improved handling of edge cases, though with some interesting trade-offs. While the weighted averages remained consistently high, the macro averages suggest that the model's performance across different attack categories became more balanced after optimization.

A particularly noteworthy observation is the model's behavior with imbalanced classes. The optimization process helped mitigate some of these challenges, but the fundamental characteristics of SVM still show through in its slightly higher sensitivity to class imbalance compared to Random Forest.

4.7 Conclusion

Random Forest and SVM both demonstrated strong performance in this experiment. Random Forest's ensemble nature provided robustness and generalization, making it the best-performing model overall. SVM, particularly with the linear kernel, offered comparable results but required more tuning to handle imbalanced classes effectively.

4.7.1 Comparative Model Analysis

Both models demonstrated exceptional performance, but with distinct characteristics:

- Random Forest showed more balanced performance across classes
- SVM demonstrated slightly higher sensitivity to class imbalance
- Both models struggled with the Impact category
- SVM showed competitive performance in major attack categories but with slightly higher FP rates

These findings have significant implications for real-world security applications, suggesting that the choice between Random Forest and SVM might depend more on specific use-case requirements than overall performance metrics alone. The analysis also highlights the importance of continuous model optimization in security contexts, where the nature of attacks constantly evolves.

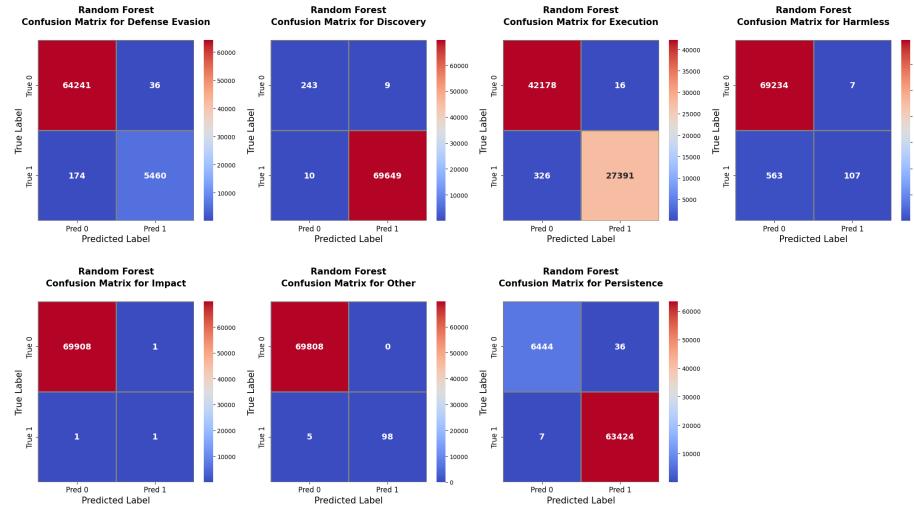


Fig. 12. Random Forest Confusion Matrices

PERFORMANCE ON TRAIN SET: Random Forest						
Model	Set	Attack	Precision	Recall	F1-Score	Accuracy
Random Forest Train	Defense Evasion		0.996432	0.986771	0.991544	0.997480
Random Forest Train	Discovery		0.991428	0.995264	0.993338	0.999896
Random Forest Train	Execution		0.996575	0.994876	0.995707	0.995887
Random Forest Train	Harmless		0.963371	0.592712	0.652848	0.992208
Random Forest Train	Impact		0.999991	0.940000	0.968081	0.999982
Random Forest Train	Other		0.999991	0.993304	0.996625	0.999982
Random Forest Train	Persistence		0.999491	0.997905	0.998697	0.999559

PERFORMANCE ON TEST SET: Random Forest						
Model	Set	Attack	Precision	Recall	F1-Score	Accuracy
Random Forest Test	Defense Evasion		0.995374	0.984278	0.989750	0.996996
Random Forest Test	Discovery		0.980173	0.982071	0.981120	0.999728
Random Forest Test	Execution		0.995873	0.993930	0.994879	0.995108
Random Forest Test	Harmless		0.965265	0.579800	0.634430	0.991847
Random Forest Test	Impact		0.749993	0.749993	0.749993	0.999971
Random Forest Test	Other		0.999964	0.975728	0.987544	0.999928
Random Forest Test	Persistence		0.999174	0.997167	0.998168	0.999385

Fig. 13. Random Forest Evaluation Metrics

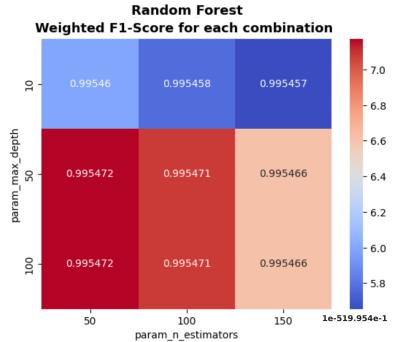


Fig. 14. Random Forest Weighted F1 Scores

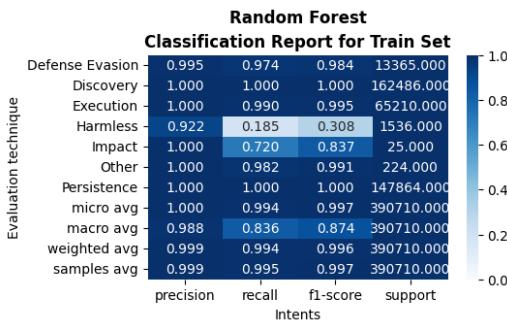


Fig. 15. Random Forest Classification Report Train Set

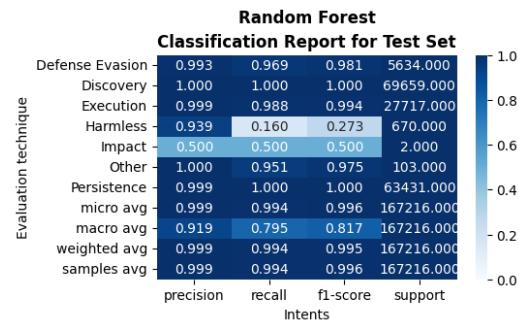


Fig. 16. Random Forest Classification Report Test Set

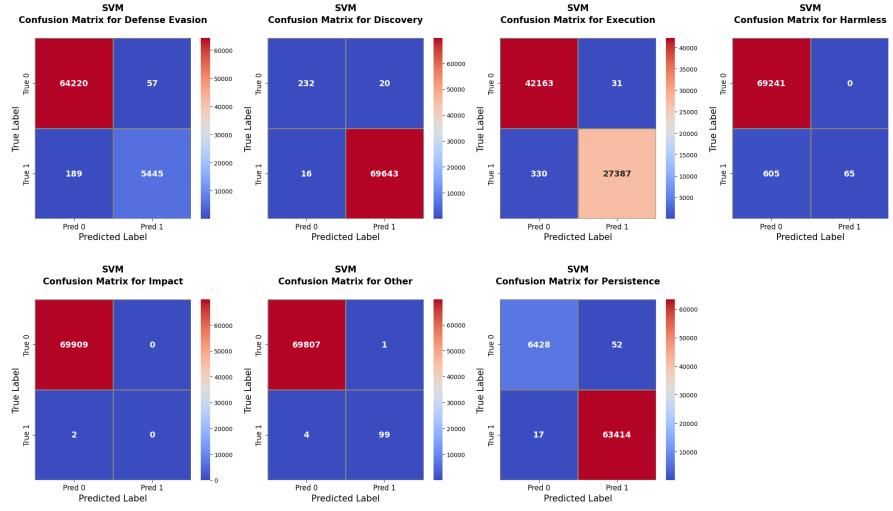


Fig. 17. SVM Confusion Matrices

PERFORMANCE ON TRAIN SET: SVM					
Model Set	Attack	Precision	Recall	F1-Score	Accuracy
SVM Train	Defense Evasion	0.993578	0.984850	0.989167	0.996769
SVM Train	Discovery	0.958840	0.964570	0.961687	0.999399
SVM Train	Execution	0.995982	0.994316	0.995131	0.995335
SVM Train	Harmless	0.995219	0.545895	0.581878	0.991442
SVM Train	Impact	0.499923	0.500000	0.499962	0.999847
SVM Train	Other	0.999966	0.975446	0.987397	0.999933
SVM Train	Persistence	0.997561	0.997123	0.997342	0.999099
PERFORMANCE ON TEST SET: SVM					
Model Set	Attack	Precision	Recall	F1-Score	Accuracy
SVM Test	Defense Evasion	0.993353	0.982783	0.987999	0.996481
SVM Test	Discovery	0.967598	0.960203	0.963871	0.999485
SVM Test	Execution	0.995552	0.993680	0.994459	0.994836
SVM Test	Harmless	0.995669	0.548507	0.586260	0.991346
SVM Test	Impact	0.499986	0.500000	0.499993	0.999971
SVM Test	Other	0.994971	0.980575	0.987667	0.999928
SVM Test	Persistence	0.998271	0.995854	0.997059	0.999013

Fig. 18. SVM Evaluation Metrics

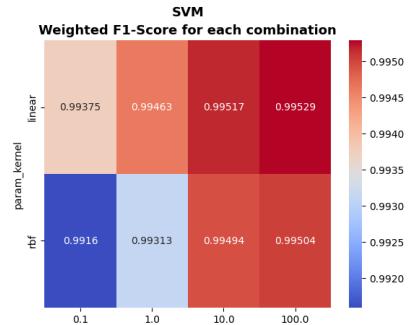


Fig. 19. SVM Weighted F1 Scores

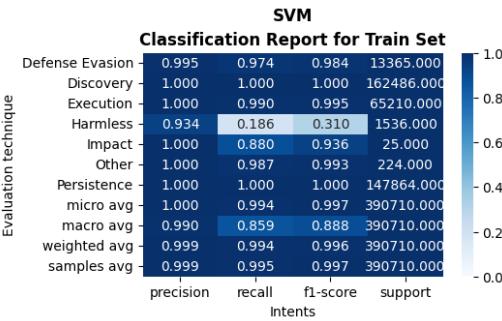


Fig. 20. SVM Classification Report Train Set

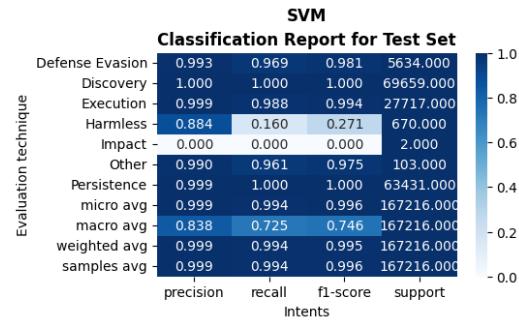


Fig. 21. SVM Classification Report Test Set

5 UNSUPERVISED LEARNING - CLUSTERING

5.1 Introduction

Unsupervised learning, a powerful branch of machine learning, was applied in this project to gain insights from SSH attack data. The primary focus was on leveraging clustering methods to group similar attack sessions based on their intrinsic patterns and characteristics. By analyzing these groups, the study aimed to uncover hidden relationships and categorize different attack intents and behaviors without relying on predefined labels.

5.2 Data Preparation

The dataset chosen was the one generated through the TF-IDF vectorization technique. This was made because it was essential to start with a dataset that represented in the best way the frequency and the importance of words, making each word as a dimension of our vector.

5.3 Clustering Methods

Clustering techniques were employed to uncover natural groupings within the dataset, providing insights into SSH attack patterns. The following methods were used:

- **K-Means Clustering:** The algorithm iteratively assigns each data point to the nearest cluster centroid and updates the centroids until convergence. The Elbow Method was applied to determine the optimal number of clusters by examining the total within-cluster sum of squares (inertia). Silhouette scores were also calculated to evaluate the cohesion and separation of clusters, ensuring the clustering results were meaningful and well-separated.
- **Gaussian Mixture Model (GMM):** Unlike K-Means, GMM considers the probability of each data point belonging to a cluster, providing a more flexible and nuanced clustering approach. The optimal number of clusters was determined using a combination of log-likelihood scores, which measure how well the model fits the data, and silhouette analysis to validate cluster quality. This dual approach ensured the GMM provided reliable and interpretable clustering results.

5.4 Clustering Evaluation Techniques

5.4.1 K-Means Clustering

The evaluation of the K-Means clustering results is based on the Elbow Method and the Silhouette Score, which provide insights into the optimal number of clusters for the dataset.

The Elbow Method graph shows a steep decline in clustering error between 3 and 6 clusters, followed by a more gradual decrease as the number of clusters increases. The point of inflection, or "elbow," appears around 6 clusters, suggesting that adding more clusters beyond this point results in diminishing improvements in minimizing intra-cluster variance.

The Silhouette Score graph exhibits a rapid increase up to 5 clusters, reaching a stable high value of approximately 0.95. A drop is observed around 8 clusters, after which the score gradually increases again, peaking beyond 12 clusters. This pattern indicates that a smaller number of clusters (around 5-6) achieves a strong balance of cohesion and separation, while additional clusters beyond 12 continue to refine the structure with marginal improvements.

Considering both metrics, the optimal number of clusters for K-Means is likely between 5 and 6, ensuring a trade-off between clustering accuracy and computational efficiency.

5.4.2 Gaussian Mixture Model (GMM)

The evaluation of the GMM clustering results is based on the Silhouette Score and the Log-Likelihood Score, both of which provide insights into the quality of the clustering structure.

The Silhouette Score graph shows a sharp increase up to 5 clusters, reaching a stable high value around 0.95. A slight drop is observed at 8 clusters, followed by a steady increase, with the highest scores occurring beyond 12 clusters. This suggests that increasing the number of clusters generally improves separation and cohesion, though the optimal balance appears to be around 6 clusters, where the highest stable performance is first achieved.

The Log-Likelihood Score graph indicates a rapid increase from 3 to 5 clusters, after which the improvements become more gradual. Beyond 12 clusters, the score stabilizes, indicating diminishing returns in model fitting. This suggests that while increasing the number of clusters provides better representation of the data, the most significant improvements occur within the first few increments.

Considering both metrics, an optimal cluster configuration is likely between 6 and 8 clusters, balancing cluster separation, model likelihood, and computational efficiency..

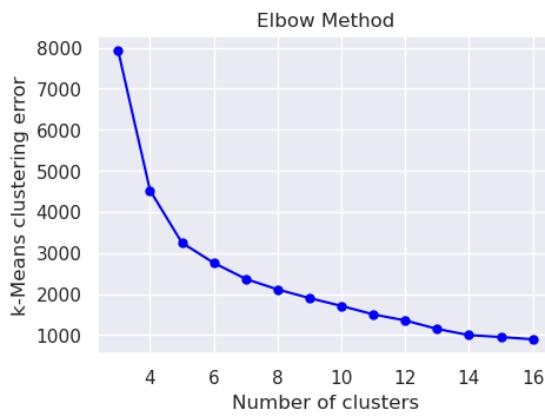


Fig. 22. K-means Elbow Method

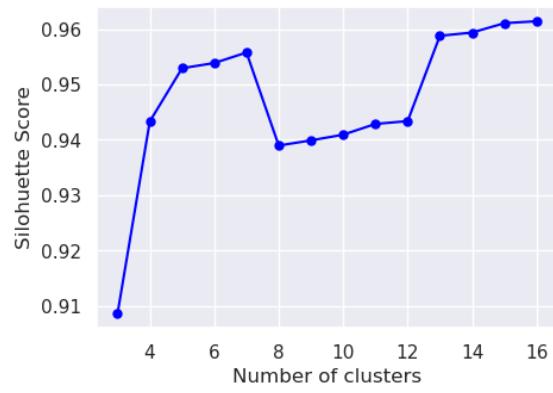


Fig. 23. K-means Silhouette Score

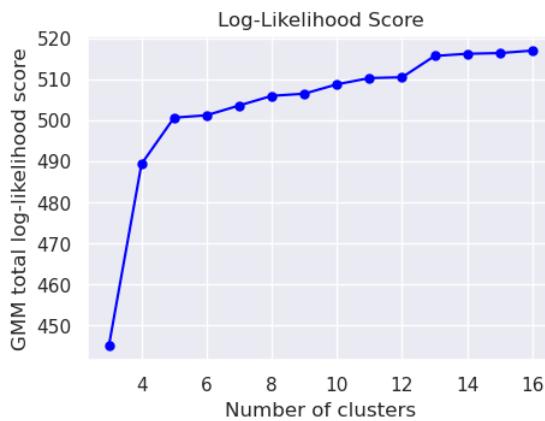


Fig. 24. GMM Log-Likelihood Score

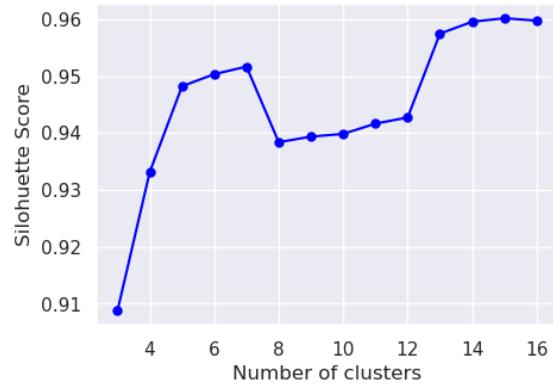


Fig. 25. GMM Silhouette Score

5.5 Hyperparameter Tuning

Hyperparameter tuning was conducted to optimize the performance of both clustering methods. For K-Means, parameters such as the initialization method (`k-means++` and `random`), the number of initializations, and the maximum number of iterations were fine-tuned using a grid search approach. Similarly, GMM parameters including the initialization method (`kmeans`), covariance type (`full` and `spherical`), and tolerance were optimized. These steps ensured the models were tailored to the dataset, resulting in better clustering outcomes.

5.6 Clusters Visualization

To visualize the clustering results, t-SNE dimensionality reduction was applied. Its functionalities makes it an excellent choice for visualizing clusters in datasets where direct interpretation is difficult due to high dimensionality. By projecting the data into a two-dimensional space, t-SNE enables us to identify patterns and groupings that may not be evident in the original feature space. The two-dimensional plots provided a clear representation of the clusters formed by both K-Means and GMM, highlighting their separability and internal consistency.

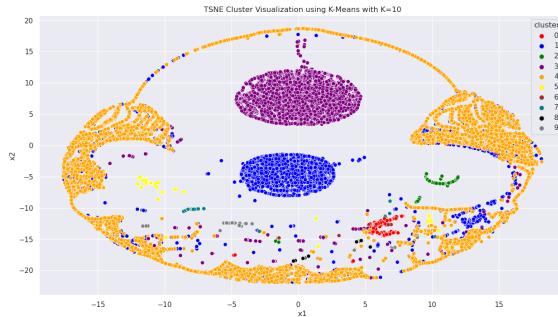


Fig. 26. t-SNE Visualization of K-Means Clusters

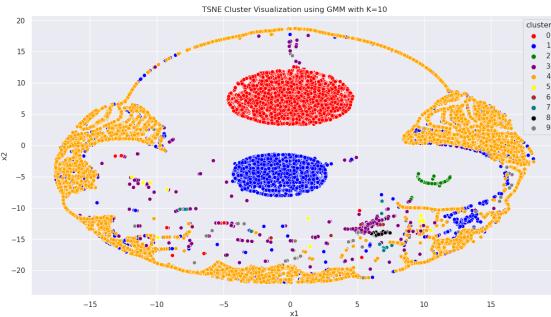


Fig. 27. t-SNE Visualization of GMM Clusters

5.6.1 K-Means Visualization

The t-SNE visualization provides a two-dimensional representation of the 10 clusters identified by the K-Means algorithm. This graph highlights the spatial distribution and separability of the clusters in the dataset.

Distinct cluster groupings are visible, with some clusters (e.g., the orange and purple clusters) forming highly compact and dense regions. This indicates strong intra-cluster similarity and effective separation from other clusters. Conversely, a few clusters (e.g., the red and yellow clusters) are more dispersed, suggesting potential overlap or variability in their features.

The visualization demonstrates that K-Means effectively partitions the dataset into meaningful groups, with well-defined boundaries for most clusters. However, the presence of dispersed clusters may reflect the complexity of certain patterns in the dataset, indicating that some attack behaviors share overlapping features. Overall, the t-SNE plot validates the clustering results by showing clear differentiation among the majority of the clusters.

5.6.2 GMM Visualization

The clusters exhibit distinct groupings, with some (e.g., the red and blue clusters) forming tightly packed regions that indicate strong intra-cluster similarity and minimal overlap. However, other clusters (e.g., the orange cluster) are more dispersed, reflecting the probabilistic nature of GMM, which allows for overlapping data points and accounts for uncertainties in cluster assignment.

Compared to the K-Means visualization, the GMM-based clusters appear more balanced in size and density. This suggests that GMM is effectively capturing the natural variability and overlapping characteristics within

the dataset. Overall, the t-SNE plot validates the clustering results, showing that GMM provides a nuanced partitioning of the dataset while accommodating the complexities inherent in the data.

5.7 Clusters Analysis

5.7.1 Word Cloud Representation

Word clouds were generated for each cluster to highlight the most significant terms. These visualizations provided an intuitive understanding of the key features within each cluster by emphasizing frequently occurring terms. The approach helped in identifying the distinguishing characteristics of each cluster, offering insights into the behavioral patterns and intents associated with different attack sessions.

The word clouds revealed dominant keywords for specific clusters, such as commands, parameters, or phrases frequently used in SSH attacks. This information serves as a valuable reference for understanding the nature of the clustered attack sessions, aiding in further analysis and interpretation of the underlying data.

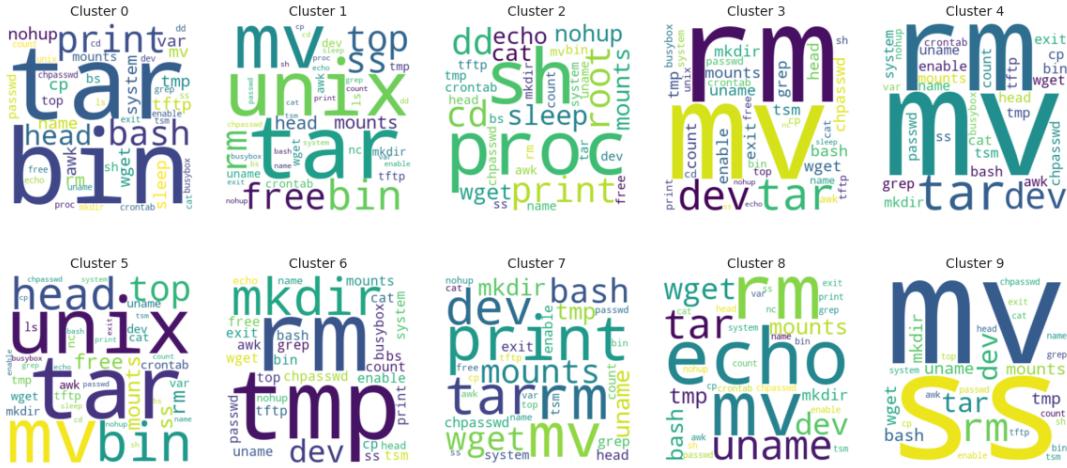


Fig. 28. Word Clouds for Each Cluster

5.7.2 Community Detection

Graph-based community detection was performed within selected clusters to identify subgroups. This approach involves constructing a graph where nodes represent features or sessions, and edges indicate relationships or co-occurrences within the data. The greedy_modularity_communities method was applied to detect communities, maximizing modularity to ensure well-defined groups.

The analysis revealed meaningful relationships and substructures within the clusters, helping to further refine the understanding of the dataset's internal patterns. These subgroups highlight finer-grained distinctions within the clusters, such as frequently occurring command sequences or behavioral traits common to specific attack types. Visualizations of these detected communities provide insights into the interconnectedness of the data and potential hierarchical structures in SSH attack patterns.

These visualizations demonstrate that distinct subgroups exist within Cluster 0. For K-Means, the graph reveals tight communities (e.g., green or pink) that suggest strong intra-community similarities, whereas more dispersed connections (e.g., cyan or orange) indicate weaker relationships. In contrast, the GMM graph shows balanced community sizes, reflecting GMM's probabilistic approach, which accounts for overlapping data points.

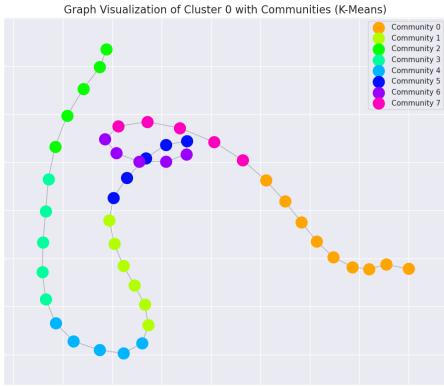


Fig. 29. Community Detection in Cluster 0 (K-Means).

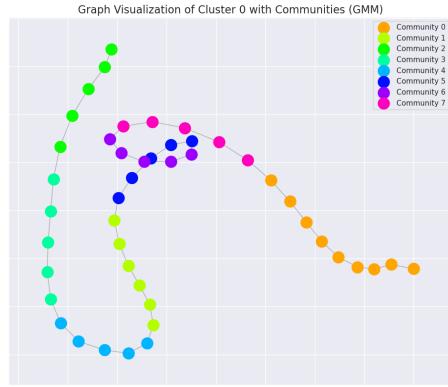


Fig. 30. Community Detection in Cluster 0 (GMM).

Overall, these detected communities provide finer-grained categorizations of attack patterns, enabling a deeper understanding of SSH attack behaviors beyond standard clustering methods.

5.8 Conclusion

This analysis provided valuable insights into the patterns of SSH attacks through clustering. The K-Means and GMM algorithms both effectively identified meaningful clusters, as supported by validation metrics and visualizations. Hyperparameter tuning further enhanced the performance of both models. The results demonstrate the potential of unsupervised learning in uncovering hidden patterns in complex datasets, providing a foundation for future applications such as anomaly detection and improved cybersecurity strategies.

6 LANGUAGE MODEL EXPLORATION

6.1 Introduction

This section documents the steps taken to fine-tune a pre-trained BERT model for multi-label classification. The focus was on customizing the final classification layer and training the model on the `Set_Fingerprint` intents. Performance metrics and visualizations are presented to evaluate the effectiveness of the model.

6.2 Data Preprocessing

The dataset was processed as follows:

- The `full_session` column, which contains lists of words, was filtered to retain only words that appeared at least in 1% of all the sessions. This step reduced the vocabulary to frequent words only.
- The `Set_Fingerprint` column was preprocessed for multi-label encoding using the `MultiLabelBinarizer` class, enabling the representation of intents as binary vectors.
- The dataset was split into 60% training, 20% validation, and 20% testing sets.

6.3 Tokenization

The tokenization process is a critical step in preparing text data for input into a BERT model. We utilized the pre-trained BERT tokenizer, specifically designed for the `bert-base-uncased` model, to convert our text data into input IDs and attention masks. Using the pre-trained tokenizer ensures consistency with BERT's pre-training,

maintaining representation integrity and optimizing model performance. These strategies enhance efficiency, compatibility, and overall performance in the text processing pipeline.

6.4 Model Architecture

A pre-trained `bert-base-uncased` model was used, with a custom linear classification head added to predict the intents. The classifier had an input size of 768 (BERT's hidden layer dimensions) and an output size equal to the number of intents. The model architecture is as follows:

- `CustomBERTModel` : A custom model class that inherits from `torch.nn.Module`.
- `BertModel` : The pre-trained BERT model.
- `classifier` : A dense layer with a sigmoid activation function that maps the BERT output to the number of intents.

The choice of a dense layer with sigmoid activation is motivated by the multi-label nature of our classification task. In multi-label classification, each instance can be associated with multiple labels simultaneously. The sigmoid activation allows each output neuron to independently predict the probability of its corresponding intent being present, which is essential for handling non-mutually exclusive labels.

6.5 Training Process

- The model was trained using the AdamW optimizer with a learning rate of 4×10^{-5} , betas of (0.9, 0.98), and an epsilon of 1×10^{-6} .
- A linear learning rate scheduler was used with no warmup steps and a total of `num_training_steps` steps.
- The `BCEWithLogitsLoss` loss function was applied for multi-label classification.
- The model was fine-tuned for 4 epochs, with a batch size of 16. Specifically, only the last layer (the custom classification head) was trained, while the pre-trained BERT layers were kept frozen.

6.6 Evaluation Metrics

- Metrics included precision, recall, F1-score, and ROC-AUC for each intent class.
- ROC curves were generated for each class to visualize the trade-off between true positive rate and false positive rate.
- Training and validation loss curves were plotted to monitor the model's learning progress over epochs.

6.7 Results

6.7.1 Training and Validation Loss

The training and validation loss curves shown in Figure 31 indicate a consistent decrease in both losses over the epochs, suggesting that the model is learning effectively. The training loss and validation loss both trend downwards, with no significant signs of overfitting, as the validation loss does not increase while the training loss decreases. This convergence implies that the model generalizes well to the validation data.

However, it is important to note that while the losses show a consistent decrease, we do not know for how much longer this trend could continue. The optimal number of epochs remains unclear, as we stopped training after 4 epochs due to resource constraints. Further experimentation with additional epochs could provide deeper insights into the model's learning dynamics and potentially improve performance.

6.7.2 ROC Curves

The ROC curves for each class, as depicted in Figure 32, show excellent performance for most classes, with AUC



Fig. 31. Training and Validation Loss

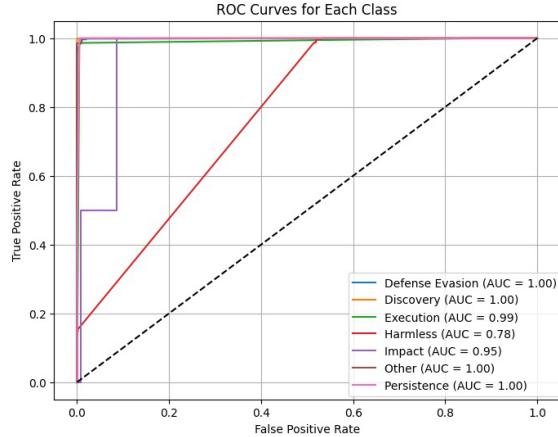


Fig. 32. ROC Curves for Each Class

values close to 1.00 for Defense Evasion, Discovery, Execution, Other, and Persistence. Notably, the Harmless class has a lower AUC of 0.78, indicating potential difficulties in distinguishing this class from others. This could suggest class imbalance or other challenges specific to identifying harmless intents.

6.7.3 Test Metrics

The model achieved the following high weighted metrics on the test set:

- Precision: 0.9974
- Recall: 0.9927
- F1-Score: 0.9937
- ROC-AUC: 0.9971

These weighted metrics indicate superior performance, capturing most positive cases accurately while accounting for class imbalance.

6.8 Conclusion

The BERT-based model demonstrated exceptional performance in classifying SSH attack sessions into MITRE ATT&CK tactics, achieving high precision, recall, F1-score, and ROC-AUC values. Notably, the model excelled in identifying malicious intents, as evidenced by the high AUC values for classes such as Defense Evasion and Persistence. However, the "Harmless" class presented challenges, indicating potential issues with class imbalance and feature representation.

Several factors contribute to these findings. The limited number of examples for harmless activities may have hindered the model's ability to learn distinguishing features effectively. Future research could explore fine-tuning more layers of the BERT model to enhance performance, or adding more instances and sessions of the lesser classes, making the dataset more balanced. Additionally, feature engineering and exploring different architectures could improve the model's discriminative power.

In summary, this study contributes to the field by successfully applying a BERT-based model to classify SSH attack sessions, achieving high performance, and identifying areas for improvement. The insights gained suggest promising directions for enhancing the model's capabilities in distinguishing harmful from harmless intents, thereby advancing cybersecurity defenses.

7 CONCLUSION

7.1 Summary of Key Findings

In this project, we explored various techniques for analyzing and classifying SSH shell attack logs. The primary objectives were to preprocess the data, perform exploratory data analysis, implement supervised and unsupervised learning models, and leverage advanced language models for classification tasks. Here, we summarize the key findings from each section of the project.

Data Exploration and Pre-processing: We began by loading and inspecting the dataset, identifying missing values, and handling duplicates. Temporal analysis revealed significant variations in attack frequencies over time, with notable peaks during specific hours and months. Feature extraction and common words analysis provided insights into the most frequent commands and intents used in the attack sessions.

Supervised Learning - Classification: We implemented and evaluated several machine learning models, including Logistic Regression, Random Forest, and Support Vector Machine (SVM). Hyperparameter tuning improved the performance of these models, and the result analysis highlighted the strengths and weaknesses of each approach.

Unsupervised Learning - Clustering: Clustering techniques, such as K-Means and Gaussian Mixture Models (GMM), were used to group similar attack sessions. The elbow method and silhouette analysis helped determine the optimal number of clusters. Cluster visualization using t-SNE provided a clear representation of the clusters, and cluster analysis revealed common patterns and behaviors within each group.

Language Model Exploration: We explored the use of advanced language models, such as BERT, for classifying attack session tactics. Fine-tuning the pretrained BERT model on our dataset improved classification performance. Learning curves indicated the optimal number of epochs for training, helping to avoid overfitting.

7.2 Challenges Faced

Throughout the project, we encountered several challenges that required careful consideration and problem-solving.

Data Quality and Preprocessing: Handling missing values, duplicates, and inconsistencies in the dataset was a critical step. Ensuring the data was clean and well-prepared for analysis required significant effort. Additionally, the unstructured nature of the session text posed challenges for text representation and feature extraction.

Model Selection and Tuning: Selecting appropriate machine learning models and tuning their hyperparameters was a complex task. Balancing model complexity with performance and avoiding overfitting required iterative experimentation and validation.

Computational Resources: Training advanced language models, such as BERT, required substantial computational resources. Efficiently managing these resources and optimizing the training process was essential to achieve timely results.

Interpretability of Results: Interpreting the results of clustering and classification models, especially in the context of cybersecurity, was challenging. Ensuring that the findings were meaningful and actionable required careful analysis and domain knowledge.

7.3 Future Work

Based on the findings and challenges encountered in this project, we propose several directions for future work.

Enhanced Feature Engineering: Further exploration of feature engineering techniques, such as incorporating domain-specific knowledge and using advanced text representation methods, could improve model performance. Experimenting with additional features, such as network metadata and contextual information, may provide deeper insights into attack patterns.

Advanced Model Architectures: Exploring more advanced model architectures, such as transformer-based models and deep neural networks, could enhance classification accuracy. Transfer learning with other pretrained models and ensemble methods may also yield better results.

Real-time Analysis and Detection: Implementing real-time analysis and detection systems for SSH shell attacks could provide immediate insights and responses to potential threats. Integrating the models developed in this project into a real-time monitoring framework would be a valuable extension.

Broader Dataset and Generalization: Expanding the dataset to include a wider range of attack types and sources would improve the generalizability of the models. Collaborating with other organizations to share data and insights could enhance the robustness and applicability of the findings.

7.4 Conclusion

This project demonstrated the potential of machine learning and advanced language models for analyzing and classifying SSH shell attack logs. By leveraging various techniques, we gained valuable insights into attack patterns and behaviors, which can inform cybersecurity strategies and defenses. Despite the challenges faced, the results highlight the importance of data-driven approaches in enhancing cybersecurity threat detection and response capabilities. Future work in this area holds promise for further advancements and practical applications in the field of cybersecurity.